

In [ ]:

```

1 from __future__ import print_function, division
2 import matplotlib.pyplot as plt
3 import numpy as np
4 # import tensorflow as tf
5 # import tensorflow.keras as keras
6 import torch
7 import torch.nn as nn
8 import torch.optim as optim
9 from torch.optim import lr_scheduler
10 import numpy as np
11 import torchvision
12
13 from torchvision import datasets, models, transforms
14 import matplotlib.pyplot as plt
15 import time
16 import os
17 import copy
18
19 # plt.ion() # interactive mode

```

In [ ]:

```

1 data_transforms = {
2     'train': transforms.Compose([
3         # transforms.RandomResizedCrop(224),
4         # transforms.RandomHorizontalFlip(),
5         transforms.ToTensor(),
6         transforms.Normalize([0.4914, 0.4822, 0.4465], [0.2023, 0.1994, 0.2023]),
7     ]),
8     'val': transforms.Compose([
9         # transforms.Resize(256),
10        # transforms.CenterCrop(224),
11        transforms.ToTensor(),
12        transforms.Normalize([0.4914, 0.4822, 0.4465], [0.2023, 0.1994, 0.2023]),
13    ]),
14 }
15
16
17
18 image_datasets = {'train': datasets.CIFAR100('data', train=True,
19                                             transform=data_transforms['train']),
20                  'val': datasets.CIFAR100('data', train=False,
21                                           transform=data_transforms['val'])}
22 dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=32,
23                                             shuffle=True) for x in ['train', 'val']}
24 dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
25 class_names = image_datasets['train'].classes
26
27 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

```

Files already downloaded and verified

Files already downloaded and verified

CIFAR-100 Dataset:


No of classes = 100 \ No of images from each class = 600 \


```

In [ ]: 1 def imshow(inp, title=None):
        2     """Imshow for Tensor."""
        3     inp = inp.numpy().transpose((1, 2, 0))
        4     mean = np.array([0.485, 0.456, 0.406])
        5     std = np.array([0.229, 0.224, 0.225])
        6     inp = std * inp + mean
        7     inp = np.clip(inp, 0, 1)
        8     plt.imshow(inp)
        9     if title is not None:
       10         plt.title(title)
       11     plt.pause(0.001)
       12
       13
       14     appind=[]
       15     sqind = []
       16     while(len(appind)<2 or len(sqind)<2):
       17         inputs, classes = next(iter(dataloaders['train']))
       18         for i in range(len(classes)):
       19             if(len(appind)==2 and len(sqind)==2):
       20                 break
       21             if(classes[i]==1 and len(appind)<2):
       22                 appind.append(inputs[i])
       23             if(classes[i]==9 and len(sqind)<2):
       24                 sqind.append(inputs[i])
       25     print(class_names[1])
       26     imshow(appind[0])
       27     imshow(appind[1])
       28     plt.show()
       29     print(class_names[9])
       30     imshow(sqind[0])
       31     plt.show()
       32     imshow(sqind[1])
       33     plt.show()
       34     print(len(dataloaders['train']),len(dataloaders['val']))

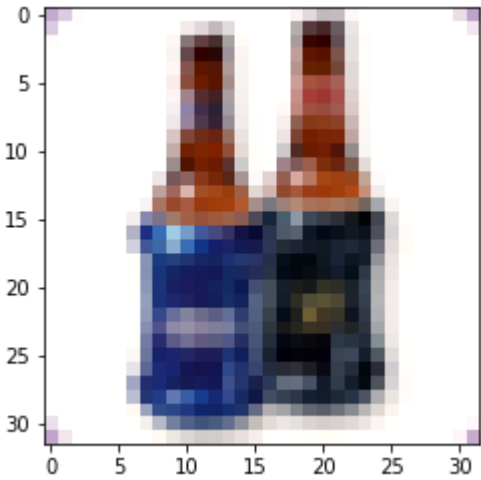

```

aquarium\_fish





bottle



782 157

```

In [ ]: 1 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
        2     since = time.time()
        3
        4     best_model_wts = copy.deepcopy(model.state_dict())
        5     best_acc = 0.0
        6     acc = []
        7     for epoch in range(num_epochs):
        8         # print('Epoch {}/{}'.format(epoch, num_epochs - 1))
        9         # print('-' * 10)
       10         epoch_time = time.time()
       11         # Each epoch has a training and validation phase
       12         for phase in ['train', 'val']:
       13             if phase == 'train':
       14                 model.train() # Set model to training mode
       15             else:
       16                 model.eval() # Set model to evaluate mode
       17
       18             running_loss = 0.0
       19             running_corrects = 0
       20
       21             # Iterate over data.
       22             for inputs, labels in dataloaders[phase]:
       23                 inputs = inputs.to(device)
       24                 labels = labels.to(device)
       25
       26                 # zero the parameter gradients
       27                 optimizer.zero_grad()
       28
       29                 # forward
       30                 # track history if only in train
       31                 with torch.set_grad_enabled(phase == 'train'):
       32                     outputs = model(inputs)
       33                     _, preds = torch.max(outputs, 1)
       34                     loss = criterion(outputs, labels)
       35
       36                 # backward + optimize only if in training phase
       37                 if phase == 'train':
       38                     loss.backward()
       39                     optimizer.step()
       40
       41                 # statistics
       42                 running_loss += loss.item() * inputs.size(0)
       43                 running_corrects += torch.sum(preds == labels.data)
       44             if phase == 'train':
       45                 scheduler.step()
       46
       47             epoch_loss = running_loss / dataset_sizes[phase]
       48             epoch_acc = running_corrects.double() / dataset_sizes[phase]
       49
       50             print('{} Loss: {:.4f} Acc: {:.4f}'.format(
       51                 phase, epoch_loss, epoch_acc))
       52
       53             # deep copy the model
       54             if phase == 'val' and epoch_acc > best_acc:
       55                 best_acc = epoch_acc
       56                 best_model_wts = copy.deepcopy(model.state_dict())

```

```

57         acc.append(epoch_acc.item())
58     #         print(acc)
59     #     if(epoch>10):
60     #         if(np.mean(acc[-10:])>epoch_acc):
61     #             break
62
63     # if(epoch>10):
64     #     if(np.mean(acc[-10:])>epoch_acc):
65     #         break
66     print('Epoch time ',time.time()-epoch_time)
67     print()
68
69     time_elapsed = time.time() - since
70     print('Training complete in {:.0f}m {:.0f}s'.format(
71         time_elapsed // 60, time_elapsed % 60))
72     print('Best val Acc: {:.4f}'.format(best_acc))
73
74     # load best model weights
75     model.load_state_dict(best_model_wts)
76     return model

```

In [ ]:

```

1 def visualize_model(model, num_images=6):
2     was_training = model.training
3     model.eval()
4     images_so_far = 0
5     fig = plt.figure()
6
7     with torch.no_grad():
8         for i, (inputs, labels) in enumerate(dataloaders['val']):
9             inputs = inputs.to(device)
10            labels = labels.to(device)
11
12            outputs = model(inputs)
13            _, preds = torch.max(outputs, 1)
14
15            for j in range(inputs.size()[0]):
16                images_so_far += 1
17                ax = plt.subplot(num_images//2, 2, images_so_far)
18                ax.axis('off')
19                ax.set_title('predicted: {}'.format(class_names[preds[j]]))
20                imshow(inputs.cpu().data[j])
21
22            if images_so_far == num_images:
23                model.train(mode=was_training)
24                return
25     model.train(mode=was_training)

```

```
In [ ]: 1 # 1.b
2 model_ft = models.resnet50(pretrained=True)
3 num_fts = model_ft.fc.in_features
4 # Here the size of each output sample is set to 2.
5 # Alternatively, it can be generalized to nn.Linear(num_fts, len(class_name
6 model_ft.fc = nn.Linear(num_fts, 100)
7
8 model_ft = model_ft.to(device)
9
10 criterion = nn.CrossEntropyLoss()
11
12 # Observe that all parameters are being optimized
13 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
14
15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=20, gamma=0.1
```

```
In [ ]: 1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
2                               num_epochs=60)
```

train Loss: 0.1670 Acc: 0.9500  
val Loss: 1.9103 Acc: 0.5982  
Epoch time 58.50404477119446

train Loss: 0.1301 Acc: 0.9613  
val Loss: 1.9514 Acc: 0.5944  
Epoch time 57.047507524490356

train Loss: 0.1153 Acc: 0.9658  
val Loss: 1.9796 Acc: 0.6012  
Epoch time 57.02501702308655

train Loss: 0.0951 Acc: 0.9720  
val Loss: 2.0204 Acc: 0.5990  
Epoch time 57.272584438323975

train Loss: 0.0932 Acc: 0.9724  
val Loss: 2.0502 Acc: 0.5947  
Epoch time 57.72999143600464

```
In [ ]: 1
```

```

In [ ]: 1 # 1.c
        2 model_ft = models.resnet50(pretrained=True)
        3 num_fts = model_ft.fc.in_features
        4
        5 model_ft.fc = nn.Linear(num_fts, 100)
        6
        7 model_ft = model_ft.to(device)
        8
        9 criterion = nn.CrossEntropyLoss()
        10
        11 # Observe that all parameters are being optimized
        12 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.01, momentum=0.9)
        13
        14 # Choosing large step_size so that the learning rate is not changed
        15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=10000, gamma=

```

```

In [ ]: 1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
        2                             num_epochs=200)

```

```
train Loss: 0.0055 Acc: 0.9988
```

```
val Loss: 2.8747 Acc: 0.5826
```

```
Epoch time 58.596367597579956
```

```
train Loss: 0.0064 Acc: 0.9983
```

```
val Loss: 2.8930 Acc: 0.5801
```

```
Epoch time 57.54751014709473
```

```
train Loss: 0.0054 Acc: 0.9985
```

```
val Loss: 3.0666 Acc: 0.5672
```

```
Epoch time 56.887218713760376
```

```
train Loss: 0.0052 Acc: 0.9985
```

```
val Loss: 2.9798 Acc: 0.5787
```

```
Epoch time 56.546945571899414
```

```
train Loss: 0.0042 Acc: 0.9990
```

```
val Loss: 2.9950 Acc: 0.5794
```

```
Epoch time 57.32757496833801
```

```
In [ ]: 1
```

```
In [ ]: 1
```



```
In [ ]: 1 model_ft = models.resnet50(pretrained=True)
2 num_fts = model_ft.fc.in_features
3 # Here the size of each output sample is set to 2.
4 # Alternatively, it can be generalized to nn.Linear(num_fts, len(class_name
5 model_ft.fc = nn.Linear(num_fts, 100)
6
7 model_ft = model_ft.to(device)
8
9 criterion = nn.CrossEntropyLoss()
10
11 # Observe that all parameters are being optimized
12 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.1, momentum=0.9)
13
14 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=10000, gamma=
```

```
In [ ]: 1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
2                               num_epochs=200)
```

```
Epoch time 31.5592214104755
```

```
train Loss: 0.0745 Acc: 0.9752
val Loss: 5.7817 Acc: 0.3835
Epoch time 30.57904863357544
```

```
train Loss: 0.0691 Acc: 0.9780
val Loss: 5.8785 Acc: 0.3758
Epoch time 30.87091326713562
```

```
train Loss: 0.0616 Acc: 0.9800
val Loss: 5.7691 Acc: 0.3872
Epoch time 30.638628244400024
```

```
train Loss: 0.0663 Acc: 0.9787
val Loss: 5.8226 Acc: 0.3886
Epoch time 31.256396532058716
```

```
Training complete in 25m 52s
Best val Acc: 0.388600
```

All the three learning rates give very high training accuracies. \ The first learning rate used gives the highest validation accuracy on the target set. Validation acc = 0.6231 \ The second learning rate gives a lower accuracy. Validation acc = 0.5794 \ And the third one gives the least among the three. Validation acc = 0.0.3886 \ Therefore the learning rate which changes from 0.0001 to 0.000001 gives the best accuracy

```
In [ ]:
```

```
1
```

```
In [ ]:
```

```
1
```

```

In [ ]: 1 # 2
        2 model_conv = torchvision.models.resnet50(pretrained=True)
        3 for param in model_conv.parameters():
        4     param.requires_grad = False
        5
        6 # Parameters of newly constructed modules have requires_grad=True by default
        7 num_ftrs = model_conv.fc.in_features
        8 model_conv.fc = nn.Linear(num_ftrs, 100)
        9
        10 model_conv = model_conv.to(device)
        11
        12 criterion = nn.CrossEntropyLoss()
        13

```

```

In [ ]: 1

```

```

In [ ]: 1 # 2.a
        2 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=1, momentum=0.9)
        3
        4 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=10000, gamma
        5 model_conv = train_model(model_conv, criterion, optimizer_conv,
        6                             exp_lr_scheduler, num_epochs=200)

```

```

train Loss: 142.7233 Acc: 0.1478
val Loss: 171.0052 Acc: 0.1797
Epoch time 18.774144172668457

```

```

train Loss: 141.3311 Acc: 0.1995
val Loss: 151.4047 Acc: 0.2069
Epoch time 19.278165340423584

```

```

train Loss: 138.1913 Acc: 0.2194
val Loss: 159.6138 Acc: 0.2148
Epoch time 23.403469800949097

```

```

train Loss: 136.6077 Acc: 0.2316
val Loss: 171.3207 Acc: 0.2133
Epoch time 19.571969509124756

```

```

train Loss: 136.5718 Acc: 0.2446
val Loss: 160.7145 Acc: 0.2236
Epoch time 42.624640703201294

```

```

In [ ]: 1 model_conv = torchvision.models.resnet50(pretrained=True)
        2 for param in model_conv.parameters():
        3     param.requires_grad = False
        4
        5 # Parameters of newly constructed modules have requires_grad=True by default
        6 num_fts = model_conv.fc.in_features
        7 model_conv.fc = nn.Linear(num_fts, 100)
        8
        9 model_conv = model_conv.to(device)
       10
       11 criterion = nn.CrossEntropyLoss()
       12 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.1, momentum=0.9)
       13
       14 # Decay LR by a factor of 0.1 every 7 epochs
       15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=10000, gamma=0.1)
       16 model_conv = train_model(model_conv, criterion, optimizer_conv,
       17                           exp_lr_scheduler, num_epochs=200)

```

Epoch time 18.200000047683716

train Loss: 12.8065 Acc: 0.3433

val Loss: 18.7889 Acc: 0.2570

Epoch time 18.594185829162598

train Loss: 12.8265 Acc: 0.3414

val Loss: 20.6364 Acc: 0.2526

Epoch time 18.247527360916138

train Loss: 12.9073 Acc: 0.3421

val Loss: 19.8451 Acc: 0.2527

Epoch time 18.78353762626648

train Loss: 12.7418 Acc: 0.3437

val Loss: 18.2947 Acc: 0.2527

Epoch time 18.217689990997314

Training complete in 23m 8s

Best val Acc: 0.260000

```

In [ ]: 1 model_conv = torchvision.models.resnet50(pretrained=True)
        2 for param in model_conv.parameters():
        3     param.requires_grad = False
        4
        5 # Parameters of newly constructed modules have requires_grad=True by default
        6 num_ftrs = model_conv.fc.in_features
        7 model_conv.fc = nn.Linear(num_ftrs, 100)
        8
        9 model_conv = model_conv.to(device)
       10
       11 criterion = nn.CrossEntropyLoss()
       12 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.01, momentum=0.9)
       13
       14 # Decay LR by a factor of 0.1 every 7 epochs
       15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=10000, gamma=0.1)
       16 model_conv = train_model(model_conv, criterion, optimizer_conv,
       17                          exp_lr_scheduler, num_epochs=200)

```

Epoch time 18.62247633934021

train Loss: 2.3275 Acc: 0.4275

val Loss: 3.4606 Acc: 0.3033

Epoch time 18.950173139572144

train Loss: 2.3305 Acc: 0.4273

val Loss: 3.4212 Acc: 0.3136

Epoch time 18.578193426132202

train Loss: 2.3288 Acc: 0.4259

val Loss: 3.5612 Acc: 0.3058

Epoch time 18.40757417678833

train Loss: 2.3154 Acc: 0.4300

val Loss: 3.4103 Acc: 0.3092

Epoch time 18.22861409187317

Training complete in 15m 33s

Best val Acc: 0.317500

```
In [ ]: 1 model_conv = torchvision.models.resnet50(pretrained=True)
2 for param in model_conv.parameters():
3     param.requires_grad = False
4
5 # Parameters of newly constructed modules have requires_grad=True by default
6 num_ftrs = model_conv.fc.in_features
7 model_conv.fc = nn.Linear(num_ftrs, 100)
8
9 model_conv = model_conv.to(device)
10
11 criterion = nn.CrossEntropyLoss()
12 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.001, momentum=0.
13
14 # Decay LR by a factor of 0.1 every 7 epochs
15 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=10000, gamm
16 model_conv = train_model(model_conv, criterion, optimizer_conv,
17                           exp_lr_scheduler, num_epochs=200)
```

```
train Loss: 2.5855 Acc: 0.4072
```

```
val Loss: 2.7751 Acc: 0.3316
```

```
Epoch time 19.38201332092285
```

```
train Loss: 2.3790 Acc: 0.4015
```

```
val Loss: 2.7780 Acc: 0.3347
```

```
Epoch time 18.83073139190674
```

```
train Loss: 2.3702 Acc: 0.4085
```

```
val Loss: 2.8138 Acc: 0.3345
```

```
Epoch time 18.913022756576538
```

```
train Loss: 2.3622 Acc: 0.4093
```

```
val Loss: 2.7819 Acc: 0.3343
```

```
Epoch time 18.922974348068237
```

```
train Loss: 2.3659 Acc: 0.4080
```

```
val Loss: 2.7631 Acc: 0.3355
```

```
Epoch time 18.783432960510254
```

```
In [ ]: 1
```

2.(a) The 0.001 learning rate gives the best accuracy on the target dataset = 0.3355

2.(b) The finetuning approach gave better results than the feature extraction approach. Among all the models, the finetuning model with a small learning rate ranging from 0.0001 to 0.000001 gave the best accuracy. This is because we expect the pretrained weights to be relatively good estimates and high learning rates may distort them too quickly. The feature extractor model may have given low accuracies because the ImageNet dataset which it is trained on might not be that similar to the CIFAR-100 dataset. So the feature extraction is not as good as expected. As the resnet model has batch normalization layers, which might be input dependent, changing the input distribution drastically may affect the results.

In [ ]:

1

## Problem 2

1. In the paper which uses weakly supervised learning, they use noisy labels(hashtags) to improve the accuracy.

Whereas the semi-supervised learning paper uses a teacher-student model in which a teacher model is trained using labeled data and is used to label unseen unlabeled data.

2.(a) Yes, the models trained using hashtags are robust against noise labels. In the paper they pre-trained a ResNeXt model with 1B images and 17k labels where  $p\%$  of the labels were randomly replaced with noise.  $p = 10\%$  decreased the top-1 accuracy on ImageNet by only around 1%, and  $p = 25\%$  decreased accuracy by about 2%.

(b) Hashtags follow Zipfian distribution and it may reduce the impact of some of the classes on the overall training loss. Resampling the hashtag distribution ensures that all classes are included in training. Resampling of the hashtag distribution is important in order to obtain good transfer to ImageNet image-classification tasks. Using uniform or square-root sampling leads to an accuracy improvement of 5 to 6% irrespective of the number of ImageNet classes in the transfer task.

3.a) The goal is to improve performance using unlabeled data. This paper suggests using two models, teacher and student so that one model can benefit from the other. The teacher model is trained using labeled data. The unlabeled data is then passed through the teacher model and top-K labels from each target variable are selected from this to create a new labeled dataset. The student model is then trained on this new dataset after which it is finetuned on the previous labeled dataset. So the student model uses the labels generated from the teacher model to train itself.

Distillation is a procedure used to compress a large model into a smaller one. In distillation the teacher model makes prediction on unlabeled data, and the inferred labels are used to train the student in a supervised fashion which is very similar to the model suggested in the paper. Therefore the teacher-student model is a type of distillation technique.

The teacher and student model is needed since the teacher model selects the top-K images in the unlabeled dataset and this dataset allows us to generate a new training set. Since the student model is always 'smaller' than the teacher model, it is a distillation technique.

b)  $K$  is the number of examples that are selected from the unlabeled dataset  $U$  for each target label.  $P$  corresponds to the number of relevant classes of an image. The reason for choosing  $P > 1$  is that it is difficult to identify accurately under-represented objects, or some may be cut-off by more prominent co-occurring objects.

c) The new labeled dataset is created by selecting the top-K images from each class from the predictions of the teacher model. Yes an image from this dataset can belong to more than one class. As  $P > 1$ , if the scores of the image for more than one class lies in the top  $K$  scores of those respective classes then it will belong to all such classes.

d) Increasing K initially increases the amount of data present for the student model and thus causes the accuracy to increase. But upon increasing K further, we observe a drop in accuracy because this causes incorrect results from the output of the teacher model to be added to the new labeled data. As this increases the noise in the labels of the new dataset the accuracy decreases.

## Problem 3

1. Achieving peak FLOPS requires customized libraries with intimate knowledge of the underlying hardware. Even specially tuned libraries may fall short of peak execution by as much as 40%. Instead of trying to measure and capture every source of inefficiency in every learning framework, the paper suggests taking a small number of representative deep learning workloads which contain convolutions, pooling, dropout, and fully connected layers and run them for a short time on a single GPU. Given observed total throughput and estimated total throughput on this benchmark, fit a scaling constant to estimate a platform percent of peak (PPP) parameter which captures the average relative inefficiency of the platform compared to peak FLOPS.
2. The VGG19 model has a conv3-256 and 2 conv3-512 layers in addition to the VGG16 model. Therefore additional FLOPs =  $4,161,798,144 = 4162\text{M}$

The distribution of FLOPs for the VGG16 model is as follows-

CONV:15360M, POOL:6M, ReLU:14M, FC:124M,

Upon adding the additional FLOPs,

The CONV layers FLOPs of the VGG19 network: 19522M

And the total FLOPs of the VGG19 network is:  $15503\text{M} + 4162\text{M} = 19,665\text{M}$

So the fraction of the total FLOPs attributed by convolution layers 0.9932.

3. The measured time and sum of layerwise timing for forward pass did not match on GPUs. This is because CUDA allows asynchronous programming. Before the time is measured, an API is called to ensure that all cores have finished their tasks. This synchronization before measuring time on the GPUs results in an extra overhead. Therefore, the sum of layerwise timing on GPUs is longer than a full forward pass.

In a full forward pass, timing is only recorded at the last layer. Therefore, a core may be assigned with the computation of following layers and thus it can continuously perform the computation without synchronization. For example, after finishing the multiply-add operations for the matrix multiplication at a CONV layer, a core can continue to calculate the max function of next ReLU layer on the output of multiply-add operations. If layerwise timing is recorded, all cores have to wait until all multiply-add operations of the CONV layer have been completed.

To mitigate the overhead, they keep GPUs iteratively running the process in a way that GPU cores can continuously perform multiply-add operations without synchronization, before recording the end time. Then, the measurement overhead is amortized over all the iterations, giving accurate

timing estimates. When the number of iterations is large the measurement overhead becomes insignificant.

4. NVidia Tesla K80: double PPP = 1.87 Tflops. Forward pass on VGG requires 15503M FLOPs; as such, one forward pass on a K80 would take  $(15503 \times 10^6) / (1.87 \times 10^{12}) = 0.00829037433\text{s}$  so throughput is 120 images/sec.

GoogLeNet: Inference time:  $(1606 \times 10^6) / (1.87 \times 10^{12}) = 0.0008588\text{s/image}$ . Throughput: 1164 images/sec.

ResNet: Inference time:  $(3922 \times 10^6) / (1.87 \times 10^{12}) = 0.0020973\text{s/img}$ . Throughput: 476 images/sec.

In [ ]:

1



# Problem 4

Collaborated with Atul Balaji ab5246

In [1]:

```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import torch.nn.init as init
5 import torchvision
6 from torch.autograd import Variable
7
8 __all__ = ['resnet18', 'resnet20', 'resnet32', 'resnet44', 'resnet56', 'resnet101', 'resnet152']
9
10 def _weights_init(m):
11     classname = m.__class__.__name__
12     #print(classname)
13     if isinstance(m, nn.Linear) or isinstance(m, nn.Conv2d):
14         init.kaiming_normal_(m.weight)
15
16 class LambdaLayer(nn.Module):
17     def __init__(self, lambd):
18         super(LambdaLayer, self).__init__()
19         self.lambd = lambd
20
21     def forward(self, x):
22         return self.lambd(x)
23
24
25 class BasicBlock(nn.Module):
26     expansion = 1
27
28     def __init__(self, in_planes, planes, stride=1, option='A'):
29         super(BasicBlock, self).__init__()
30         self.conv1 = nn.Conv2d(in_planes, planes, kernel_size=3, stride=stride, padding=1)
31         self.bn1 = nn.BatchNorm2d(planes)
32         self.conv2 = nn.Conv2d(planes, planes, kernel_size=3, stride=1, padding=1)
33         self.bn2 = nn.BatchNorm2d(planes)
34
35         self.shortcut = nn.Sequential()
36         if stride != 1 or in_planes != planes:
37             if option == 'A':
38                 """
39                 For CIFAR10 ResNet paper uses option A.
40                 """
41                 self.shortcut = LambdaLayer(lambda x:
42                                             F.pad(x[:, :, ::2, ::2], (0, 0, 0, 0), 'constant', 1))
43             elif option == 'B':
44                 self.shortcut = nn.Sequential(
45                     nn.Conv2d(in_planes, self.expansion * planes, kernel_size=3, stride=stride, padding=1),
46                     nn.BatchNorm2d(self.expansion * planes)
47                 )
48
49     def forward(self, x):
50         out = F.relu(self.bn1(self.conv1(x)))
51         out = self.bn2(self.conv2(out))
52         out += self.shortcut(x)
53         out = F.relu(out)
54         return out
55
56

```

```

57 class ResNet(nn.Module):
58     def __init__(self, block, num_blocks, num_classes=10):
59         super(ResNet, self).__init__()
60         self.in_planes = 16
61
62         self.conv1 = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1, bias=True)
63         self.bn1 = nn.BatchNorm2d(16)
64         self.layer1 = self._make_layer(block, 16, num_blocks[0], stride=1)
65         self.layer2 = self._make_layer(block, 32, num_blocks[1], stride=2)
66         self.layer3 = self._make_layer(block, 64, num_blocks[2], stride=2)
67         self.linear = nn.Linear(64, num_classes)
68
69         self.apply(_weights_init)
70
71     def _make_layer(self, block, planes, num_blocks, stride):
72         strides = [stride] + [1]*(num_blocks-1)
73         layers = []
74         for stride in strides:
75             layers.append(block(self.in_planes, planes, stride))
76             self.in_planes = planes * block.expansion
77
78         return nn.Sequential(*layers)
79
80     def forward(self, x):
81         out = F.relu(self.bn1(self.conv1(x)))
82         out = self.layer1(out)
83         out = self.layer2(out)
84         out = self.layer3(out)
85         out = F.avg_pool2d(out, out.size()[3])
86         out = out.view(out.size(0), -1)
87         out = self.linear(out)
88         return out
89
90 def resnet18():
91     # return torchvision.models.resnet18()
92     return ResNet(BasicBlock, [3,3,2])
93
94 def resnet20():
95     return ResNet(BasicBlock, [3, 3, 3])
96
97
98 def resnet32():
99     return ResNet(BasicBlock, [5, 5, 5])
100
101
102 def resnet44():
103     return ResNet(BasicBlock, [7, 7, 7])
104
105
106 def resnet56():
107     return ResNet(BasicBlock, [9, 9, 9])
108
109 def resnet50():
110     return ResNet(BasicBlock, [8, 8, 8])
111
112 def test(net):
113     import numpy as np

```

```

114     total_params = 0
115
116     for x in filter(lambda p: p.requires_grad, net.parameters()):
117         total_params += np.prod(x.data.numpy().shape)
118     print("Total number of params", total_params)
119     print("Total layers", len(list(filter(lambda p: p.requires_grad and len(
120
121
122
123     for net_name in __all__:
124         if net_name.startswith('resnet'):
125             print(net_name)
126             test(globals()[net_name]())
127             print()
128

```

resnet18

Total number of params 195738

Total layers 18

resnet20

Total number of params 269722

Total layers 20

resnet32

Total number of params 464154

Total layers 32

resnet44

Total number of params 658586

Total layers 44

resnet56

Total number of params 853018

Total layers 56

resnet50

Total number of params 755802

Total layers 50

In [2]:

```

1 import torch
2 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
3 print(device)

```

cuda:0

```
In [3]: 1 import torch
2 import torchvision
3 import torchvision.transforms as transforms
4 import matplotlib.pyplot as plt
5 import numpy as np
6
7
8 transform = transforms.Compose(
9     [transforms.ToTensor(),
10      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
11
12 batch_size = 128
13
14 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
15                                         download=True, transform=transform)
16 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
17                                           shuffle=True)
18
19 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
20                                         download=True, transform=transform)
21 testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
22                                          shuffle=False)
23
24 classes = ('plane', 'car', 'bird', 'cat',
25            'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
26 def imshow(img):
27     img = img / 2 + 0.5     # unnormalize
28     npimg = img.numpy()
29     plt.imshow(np.transpose(npimg, (1, 2, 0)))
30     plt.show()
31
32
33 # get some random training images
34 dataiter = iter(trainloader)
35 images, labels = dataiter.next()
36
37 # show images
38 imshow(torchvision.utils.make_grid(images))
39 # print labels
40 print(' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
```

Files already downloaded and verified

Files already downloaded and verified



horse ship dog frog deer ship horse dog truck plane plane truck frog  
 car truck plane car plane cat dog deer ship cat car car cat de  
 er dog frog deer deer frog plane plane truck deer plane horse car hors  
 e bird frog horse ship car frog plane cat ship plane cat frog truck  
 ship dog frog ship frog frog bird car truck dog dog dog ship  
 cat deer truck car horse ship ship deer bird cat deer cat plane de  
 er car dog ship bird deer ship frog dog cat dog dog dog shi  
 p truck ship bird cat car bird dog truck ship cat truck truck car  
 truck ship horse cat dog bird horse cat plane cat deer car cat t  
 ruck deer frog truck plane frog plane dog car

In [6]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet18().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses1 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses1.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time',time.time()-s)
34 print('Finished Training')

```

```

y", line 1328, in __del__
    self._shutdown_workers()
File "/opt/conda/lib/python3.7/site-packages/torch/utils/data/dataloader.p
y", line 1320, in _shutdown_workers
    if w.is_alive():
File "/opt/conda/lib/python3.7/multiprocessing/process.py", line 151, in i
s_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
AssertionError: can only test a child process

```

```

(292, 2.1838934421539307, 0.004304968754342776)
(292, 4.227958679199219, 0.003536822247005315)
(292, 6.244589805603027, 0.007616428266355425)
(292, 8.25268292427063, 0.009005396787257751)
(292, 10.308573961257935, 0.016861269143600092)
(292, 12.375986099243164, 0.00898786271656198)
(292, 14.448589563369751, 0.006666903046663015)
epoch time 16.18087601661682

```

```
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x
```

In [20]:

```
1 textfile = open("loss1.txt", "w")
2 for element in losses1:
3     textfile.write(str(element) + "\n")
4 textfile.close()
```



In [12]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet20().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses2 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses2.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time', time.time()-s)
34 print('Finished Training')

```

```

epoch time 14.947744369506836
(348, 1.929938793182373, 0.0011401226502315768)
(348, 3.8538639545440674, 0.0011999636884940294)
(348, 5.774469614028931, 0.0010624661568281412)
(348, 7.703281402587891, 0.001675010332124954)
(348, 9.6322660446167, 0.0026279223972351805)
(348, 11.554770946502686, 0.0022304148843264853)
(348, 13.482335567474365, 0.001752686094046019)
epoch time 15.047699213027954
(349, 1.9364550113677979, 0.0016400950915674319)
(349, 3.93871808052063, 0.001300898887815752)
(349, 5.860158443450928, 0.0015589822792300802)
(349, 7.784247159957886, 0.0018144510249840096)
(349, 9.711793422698975, 0.0020217249414655474)
(349, 11.62366795539856, 0.0016196743770483502)
(349, 13.521844148635864, 0.0011128773024823631)
epoch time 15.079636335372925
(350, 1.9359312057495117, 0.0018928676866807462)
(350, 3.853344678878784, 0.0014011362253935362)
(350, 5.757847547531128, 0.001177745205124517)

```

In [21]:

```
1 textfile = open("loss2.txt", "w")
2 for element in losses2:
3     textfile.write(str(element) + "\n")
4 textfile.close()
```

In [14]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet32().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses3 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses3.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time',time.time()-s)
34 print('Finished Training')

```

```

epoch time 16.981651544570923
(348, 2.1738123893737793, 0.005050788302098078)
(348, 4.375619888305664, 0.0020687289483137237)
(348, 6.552916526794434, 0.002176484085497095)
(348, 8.717389583587646, 0.002841057912185218)
(348, 10.906378746032715, 0.004713833947993853)
(348, 13.078289031982422, 0.005461510687493909)
(348, 15.234283208847046, 0.004482235760090644)
epoch time 16.993855476379395
(349, 2.15759015083313, 0.005680434474108589)
(349, 4.330235242843628, 0.006009943694189875)
(349, 6.558507919311523, 0.007948258941000024)
(349, 8.713350534439087, 0.0064665137342300874)
(349, 10.95883822441101, 0.006662871262144146)
(349, 13.132500410079956, 0.00519731833851345)
(349, 15.293569564819336, 0.007836620545499407)
epoch time 17.066230297088623
(350, 2.193568229675293, 0.006381992861267408)
(350, 4.3531880378723145, 0.0053785650747975484)
(350, 6.49826455116272, 0.003700878180988722)

```

In [22]:

```
1 textfile = open("loss3.txt", "w")
2 for element in losses3:
3     textfile.write(str(element) + "\n")
4 textfile.close()
```

In [16]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet44().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses4 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses4.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time',time.time()-s)
34 print('Finished Training')

```

```

epoch time 19.30081868171692
(348, 2.4183807373046875, 0.0002601974845266359)
(348, 4.85472297668457, 0.00026425342429672817)
(348, 7.2866370677948, 0.0002358581078457361)
(348, 9.753231287002563, 0.0002502906684881333)
(348, 12.216313123703003, 0.00026219147143912813)
(348, 14.651654958724976, 0.0002639513215250858)
(348, 17.15526032447815, 0.0003367380144185035)
epoch time 19.15051579475403
(349, 2.4355103969573975, 0.00034760069810341074)
(349, 4.880463361740112, 0.0001450401574699327)
(349, 7.329374074935913, 0.00039748228159531174)
(349, 9.773640394210815, 0.00024213404938155708)
(349, 12.22969102859497, 0.00024203107772782274)
(349, 14.669119358062744, 0.00019463021761817353)
(349, 17.112865924835205, 0.00021934841758908933)
epoch time 19.118512868881226
(350, 2.4492437839508057, 0.00023668591358950742)
(350, 4.93298602104187, 0.0001306754963238467)
(350, 7.421787261962891, 0.00015784267174128719)

```

In [23]:

```
1 textfile = open("loss4.txt", "w")
2 for element in losses4:
3     textfile.write(str(element) + "\n")
4 textfile.close()
```

In [18]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet56().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses5 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses5.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time',time.time()-s)
34 print('Finished Training')

```

```

epoch time 21.14885187149048
(348, 2.7219815254211426, 0.00013521026914754564)
(348, 5.428853273391724, 0.000286350500517779)
(348, 8.151169300079346, 0.00011966762743826139)
(348, 10.885328531265259, 0.0001462166046436249)
(348, 13.623244762420654, 0.00028213543003273244)
(348, 16.37775468826294, 0.00025655352770367505)
(348, 19.106053590774536, 0.00015989755626592063)
epoch time 21.332329750061035
(349, 2.740628242492676, 0.0003432557903542314)
(349, 5.478779077529907, 0.00016537508560889234)
(349, 8.230273723602295, 0.00020195815500765278)
(349, 10.978095293045044, 0.00011975838684260712)
(349, 14.93930721282959, 0.00014468777816527647)
(349, 17.750129461288452, 0.00014860440536054108)
(349, 20.486787796020508, 0.00027546815077325197)
epoch time 22.712486028671265
(350, 2.745670795440674, 0.0005892272207087704)
(350, 5.470365762710571, 0.0001469808838158556)
(350, 8.200093030929565, 0.0001780625686794282)

```

```
In [1]: 1 textfile = open("loss5.txt", "w")
        2 for element in losses5:
        3     textfile.write(str(element) + "\n")
        4 textfile.close()
```

...



In [5]:

```

1 import time
2 import torch.optim as optim
3 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
4 print(device)
5 model = resnet50().to(device)
6 criterion = nn.CrossEntropyLoss()
7 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
8 losses6 = []
9 for epoch in range(350): # loop over the dataset multiple times
10     s = time.time()
11     running_loss = 0.0
12     for i, data in enumerate(trainloader, 0):
13         # print(i)
14         # get the inputs; data is a list of [inputs, labels]
15         inputs, labels = data
16         inputs, labels = inputs.to(device), labels.to(device)
17
18         # zero the parameter gradients
19         optimizer.zero_grad()
20
21         # forward + backward + optimize
22         outputs = model(inputs)
23         loss = criterion(outputs, labels)
24         loss.backward()
25         optimizer.step()
26
27         # print statistics
28         running_loss += loss.item()
29         losses6.append(loss.item())
30         if i % 50 == 49: # print every 2000 mini-batches
31             print((epoch + 1, time.time()-s, running_loss/49))
32             running_loss = 0.0
33     print('epoch time',time.time()-s)
34 print('Finished Training')

```

```

(347, 19.128620147705078, 0.0003649901712434461)
epoch time 21.215643405914307
(348, 2.61718487739563, 0.0006877457737424934)
(348, 5.144724607467651, 0.0002900304242182162)
(348, 7.704703330993652, 0.0007437641512037122)
(348, 10.253649473190308, 0.0004070300750086341)
(348, 12.803223609924316, 0.00045129853554369347)
(348, 15.361856460571289, 0.00034098081484945414)
(348, 17.948991537094116, 0.0007296348025823934)
epoch time 20.02879571914673
(349, 2.534803628921509, 0.0004356831294509383)
(349, 5.063356399536133, 0.0002361861720137127)
(349, 7.575986385345459, 0.00043401876179387374)
(349, 10.115673542022705, 0.00028960825132956844)
(349, 12.656275987625122, 0.0003594618736278047)
(349, 15.185049295425415, 0.0004238990606646272)
(349, 17.738530158996582, 0.00026988007314740774)
epoch time 19.840256690979004
(350, 2.552741289138794, 0.0006491960738625197)
(350, 5.088864088058472, 0.0012246367715510577)

```

```
In [6]: 1 textfile = open("loss6.txt", "w")
        2 for element in losses6:
        3     textfile.write(str(element) + "\n")
        4 textfile.close()
```

```
In [ ]: 1
```

In [17]:

```

1 import time
2 import torch.optim as optim
3 import torchvision.models as models
4
5 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
6 print(device)
7 model = models.resnet50().to(device)
8 criterion = nn.CrossEntropyLoss()
9 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
10 losses6 = []
11 for epoch in range(350): # loop over the dataset multiple times
12     s = time.time()
13     running_loss = 0.0
14     for i, data in enumerate(trainloader, 0):
15         # print(i)
16         # get the inputs; data is a list of [inputs, labels]
17         inputs, labels = data
18         inputs, labels = inputs.to(device), labels.to(device)
19
20         # zero the parameter gradients
21         optimizer.zero_grad()
22
23         # forward + backward + optimize
24         outputs = model(inputs)
25         loss = criterion(outputs, labels)
26         loss.backward()
27         optimizer.step()
28
29         # print statistics
30         running_loss += loss.item()
31         losses6.append(loss.item())
32         if i % 50 == 49: # print every 2000 mini-batches
33             print((epoch + 1, time.time()-s, running_loss/49))
34             running_loss = 0.0
35     train_acc = 0
36     count = 0.0
37     for i, data in enumerate(testloader, 0):
38         inputs, labels = data
39         inputs, labels = inputs.to(device), labels.to(device)
40         outputs = model(inputs)
41         _, y_pred = torch.max(outputs, 1)
42         # print(y_pred.shape, labels.shape)
43         train_acc += torch.sum(y_pred == labels)
44         count+=len(y_pred)
45     train_acc=float(train_acc/count)
46     print('Acc',train_acc)
47     if(train_acc>0.92):
48         break
49     print('epoch time',time.time()-s,train_acc/count)
50 print('Finished Training')

```

```

cuda:0
(1, 3.0936830043792725, 3.122834215358812)
(1, 6.160975694656372, 2.3637631620679582)
(1, 9.212742805480957, 2.241091256238976)

```

```
(1, 12.284221649169922, 2.1585688882944534)
(1, 15.339528799057007, 2.1069597614054776)
(1, 18.41964817047119, 2.079439844403948)
(1, 21.48569631576538, 1.9995545951687559)
Acc 0.31610000133514404
epoch time 26.452991008758545 3.1610000133514406e-05
(2, 3.066366672515869, 1.926871650072993)
(2, 6.19232702255249, 1.895579634880533)
(2, 9.281157732009888, 1.895548282837381)
(2, 12.345937490463257, 1.8633296635686134)
(2, 15.401228904724121, 1.8469602404808512)
(2, 18.446027517318726, 1.792870536142466)
(2, 21.51021122932434, 1.8003771597025346)
Acc 0.3610999882221222
epoch time 26.551684856411705 3.610999882221222e-05
```

In [5]:

```

1 import time
2 import torch.optim as optim
3 import torchvision.models as models
4
5 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
6 print(device)
7 model = resnet50().to(device)
8 criterion = nn.CrossEntropyLoss()
9 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
10 losses6 = []
11 for epoch in range(350): # loop over the dataset multiple times
12     s = time.time()
13     running_loss = 0.0
14     train_acc = 0
15     count = 0.0
16     for i, data in enumerate(trainloader, 0):
17         # print(i)
18         # get the inputs; data is a list of [inputs, labels]
19         inputs, labels = data
20         inputs, labels = inputs.to(device), labels.to(device)
21
22         # zero the parameter gradients
23         optimizer.zero_grad()
24
25         # forward + backward + optimize
26         outputs = model(inputs)
27         loss = criterion(outputs, labels)
28         loss.backward()
29         optimizer.step()
30
31         # print statistics
32         running_loss += loss.item()
33         losses6.append(loss.item())
34         if i % 50 == 49: # print every 2000 mini-batches
35             print((epoch + 1, time.time()-s, running_loss/49))
36             running_loss = 0.0
37     # train_acc = 0
38     # count = 0.0
39     # for i, data in enumerate(testloader, 0):
40     #     inputs, labels = data
41     #     inputs, labels = inputs.to(device), labels.to(device)
42     #     outputs = model(inputs)
43     #     _, y_pred = torch.max(outputs, 1)
44     #     print(y_pred.shape, labels.shape)
45     #     train_acc += torch.sum(y_pred == labels)
46     #     count+=len(y_pred)
47     train_acc=float(train_acc/count)
48     print('Acc',train_acc)
49     if(train_acc>0.92):
50         break
51     print('epoch time',time.time()-s,train_acc/count)
52 print('Finished Training')

```

```

(33, 7.665208578109741, 0.26381913770218285)
(33, 10.228975057601929, 0.28887347390457074)
(33, 12.67027759552002, 0.28931142268132193)
(33, 15.210508346557617, 0.29571513375457453)

```

```
(33, 17.721519470214844, 0.29120015428990736)
Acc 0.9042800068855286
epoch time 19.78704047203064 1.808560013771057e-05
(34, 2.5666356086730957, 0.2451854366428998)
(34, 5.6526148319244385, 0.2415621508749164)
(34, 8.839743614196777, 0.23631427695556562)
(34, 11.481876373291016, 0.2559019613022707)
(34, 14.03534746170044, 0.286735474151008)
(34, 16.591984748840332, 0.27381640094883586)
(34, 19.085593223571777, 0.27188572956591234)
Acc 0.9103999733924866
epoch time 21.150903463363647 1.820799946784973e-05
(35, 2.5165305137634277, 0.22343446770492864)

(35, 4.952503204345703, 0.21413426222849866)
(35, 7.449020147323608, 0.2245035360054094)
(35, 10.001287212880101, 0.22077179001850071)
```

In [ ]:

1

In [53]:

```
1 losses1 = []
2 textfile = open("loss1.txt", "r")
3 for element in textfile:
4     losses1.append(float(element))
5 # print(losses1)
6 textfile.close()
```

In [54]:

```
1 losses2 = []
2 textfile = open("loss2.txt", "r")
3 for element in textfile:
4     losses2.append(float(element))
5 # print(losses2)
6 textfile.close()
```

In [55]:

```
1 losses3 = []
2 textfile = open("loss3.txt", "r")
3 for element in textfile:
4     losses3.append(float(element))
5 # print(losses3)
6 textfile.close()
```

In [56]:


```
1 losses4 = []
2 textfile = open("loss4.txt", "r")
3 for element in textfile:
4     losses4.append(float(element))
5 # print(losses4)
6 textfile.close()
```

In [57]:

```
1 losses5 = []
2 textfile = open("loss5.txt", "r")
3 for element in textfile:
4     losses5.append(float(element))
5 # print(losses5)
6 textfile.close()
```

```
In [58]: 1 import matplotlib.pyplot as plt
2
3 plt.plot(range(len(losses1)),losses1)
4 plt.plot(range(len(losses2)),losses2)
5 plt.plot(range(len(losses3)),losses3)
6 plt.plot(range(len(losses4)),losses4)
7 plt.plot(range(len(losses5)),losses5)
8 plt.ylim(0,3)
```

Out[58]: (0.0, 3.0)




```
In [59]: 1 v100 = []
```



```
In [60]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses1)+1),losses1,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 v100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses1)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses1)+1),losses1,alpha=0.3,label='Loss curve')
10 plt.legend()
```

```
[ 7.01994062e-05  4.87482273e-01 -1.50206444e-01]
```


```
Out[60]: <matplotlib.legend.Legend at 0x7f9718a00b10>
```



```
In [61]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses2)+1),losses2,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 v100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses2)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses2)+1),losses2,alpha=0.3,label='Loss curve')
10 plt.legend()
```

```
[ 9.96594588e-05  4.74812073e-01 -1.20806639e-01]
```


```
Out[61]: <matplotlib.legend.Legend at 0x7f9718983e90>
```



```
In [62]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses3)+1),losses3,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 v100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses3)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses3)+1),losses3,alpha=0.3,label='Loss curve')
10 plt.legend()
```

```
[ 1.32378100e-04  4.69869572e-01 -9.88729486e-02]
```

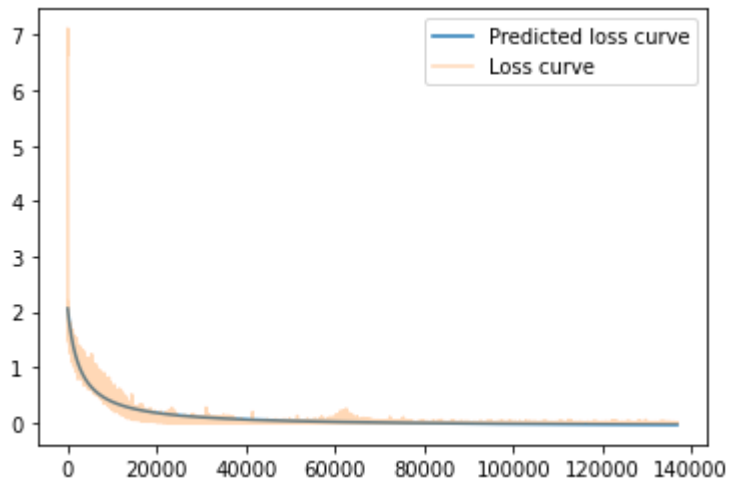
```
Out[62]: <matplotlib.legend.Legend at 0x7f97188fe9d0>
```



```
In [63]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses4)+1),losses4,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 v100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses4)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses4)+1),losses4,alpha=0.3,label='Loss curve')
10 plt.legend()
```

```
[ 1.72993949e-04  4.67240070e-01 -8.18453202e-02]
```

```
Out[63]: <matplotlib.legend.Legend at 0x7f97189d6c90>
```




```

In [64]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses5)+1),losses5,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 v100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses5)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses5)+1),losses5,alpha=0.3,label='Loss curve')
10 plt.legend()

```

```
[ 1.63617982e-04  4.65227282e-01 -8.58001223e-02]
```

Out[64]: <matplotlib.legend.Legend at 0x7f971a39a550>



```
In [65]: 1 print(v100)
```

```
[array([ 7.01994062e-05,  4.87482273e-01, -1.50206444e-01]), array([ 9.96594588
e-05,  4.74812073e-01, -1.20806639e-01]), array([ 1.32378100e-04,  4.69869572e-
01, -9.88729486e-02]), array([ 1.72993949e-04,  4.67240070e-01, -8.18453202e-0
2]), array([ 1.63617982e-04,  4.65227282e-01, -8.58001223e-02])]
```

```
In [65]: 1
```

```
In [66]: 1 import pandas as pd
          2
          3 loss_data = pd.read_pickle('merged_result_P100.pickle')
```


```
In [67]: 1 loss = []
          2 for i in loss_data:
          3     if i!='resnet50':
          4         loss.append(loss_data[i])
          5 loss.append(loss_data['resnet50'])
          6
```

```
In [68]: 1 p100 = []
          2 losses1,losses2,losses3,losses4,losses5,_ = loss
```

```
In [69]: 1 from scipy.optimize import curve_fit
          2 def func(k,beta0,beta1,beta2):
          3     return (1/(beta0*k+beta1)) + beta2
          4 popt, pcov = curve_fit(func, range(1,len(losses1)+1),losses1,p0=[0.001, 0.1,
          5 beta0,beta1,beta2 = popt
          6 p100.append(popt)
          7 print(popt)
          8 plt.plot(range(1,len(losses1)+1),[func(k,beta0,beta1,beta2) for k in range(1
          9 plt.plot(range(1,len(losses1)+1),losses1,alpha=0.3,label='Loss curve')
          10 plt.legend()
```

[1.98766092e-04 5.48874855e-01 1.76083898e-02]

Out[69]: <matplotlib.legend.Legend at 0x7f971df15210>




```

In [70]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses2)+1),losses2,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 p100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses2)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses2)+1),losses2,alpha=0.3,label='Loss curve')
10 plt.legend()

```

```
[2.90696922e-04 5.69169631e-01 5.44017628e-02]
```


```
Out[70]: <matplotlib.legend.Legend at 0x7f971dd69e10>
```



```
In [71]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses3)+1),losses3,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 p100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses3)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses3)+1),losses3,alpha=0.3,label='Loss curve')
10 plt.legend()
```

```
[3.15287961e-04 5.30339862e-01 2.07627717e-02]
```

```
Out[71]: <matplotlib.legend.Legend at 0x7f9718351dd0>
```






```

In [72]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses4)+1),losses4,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 p100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses4)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses4)+1),losses4,alpha=0.3,label='Loss curve')
10 plt.legend()

```

[3.57152914e-04 4.70425416e-01 1.78155181e-02]

Out[72]: <matplotlib.legend.Legend at 0x7f9722216950>




```

In [73]: 1 from scipy.optimize import curve_fit
2 def func(k,beta0,beta1,beta2):
3     return (1/(beta0*k+beta1)) + beta2
4 popt, pcov = curve_fit(func, range(1,len(losses5)+1),losses5,p0=[0.001, 0.1,
5 beta0,beta1,beta2 = popt
6 p100.append(popt)
7 print(popt)
8 plt.plot(range(1,len(losses5)+1),[func(k,beta0,beta1,beta2) for k in range(1
9 plt.plot(range(1,len(losses5)+1),losses5,alpha=0.3,label='Loss curve')
10 plt.legend()

```

```
[3.68076907e-04 4.27519868e-01 8.02674203e-03]
```

```
Out[73]: <matplotlib.legend.Legend at 0x7f9719726b90>
```



```

In [74]: 1 print(v100)
2 print(p100)

```

```

[array([ 7.01994062e-05,  4.87482273e-01, -1.50206444e-01]), array([ 9.96594588
e-05,  4.74812073e-01, -1.20806639e-01]), array([ 1.32378100e-04,  4.69869572e-
01, -9.88729486e-02]), array([ 1.72993949e-04,  4.67240070e-01, -8.18453202e-0
2]), array([ 1.63617982e-04,  4.65227282e-01, -8.58001223e-02])]
[array([1.98766092e-04, 5.48874855e-01, 1.76083898e-02]), array([2.90696922e-0
4, 5.69169631e-01, 5.44017628e-02]), array([3.15287961e-04, 5.30339862e-01, 2.0
7627717e-02]), array([3.57152914e-04, 4.70425416e-01, 1.78155181e-02]), array
([3.68076907e-04, 4.27519868e-01, 8.02674203e-03])]

```

```
In [74]: 1
```

In [75]:

```

1 # k80
2 import numpy as np
3
4 losses1 = np.load('resnet18_losses.npz')
5 losses2 = np.load('resnet20_losses.npz')
6 losses3 = np.load('resnet32_losses.npz')
7 losses4 = np.load('resnet44_losses.npz')
8 losses5 = np.load('resnet56_losses.npz')
9 losses1 = losses1.f.arr_0
10 losses2 = losses2.f.arr_0
11 losses3 = losses3.f.arr_0
12 losses4 = losses4.f.arr_0
13 losses5 = losses5.f.arr_0
14 # print(losses1.keys)
15 k80 = []
16
17 losses6 = np.load('resnet50_new.npz')
18 losses6 = losses6.f.a
19 print(losses6)

```

```

[7.00012493e+00 6.51588106e+00 6.20615292e+00 ... 4.14951630e-02
 3.50350663e-02 4.87744343e-03]

```

In [76]:

```

1 from scipy.optimize import curve_fit
2 import matplotlib.pyplot as plt
3
4 def func(k,beta0,beta1,beta2):
5     return (1/(beta0*k+beta1)) + beta2
6 popt, pcov = curve_fit(func, range(1,len(losses1)+1),losses1,p0=[0.001, 0.1,
7 beta0,beta1,beta2 = popt
8 k80.append(popt)
9 print(popt)
10 plt.plot(range(1,len(losses1)+1),[func(k,beta0,beta1,beta2) for k in range(1
11 plt.plot(range(1,len(losses1)+1),losses1,alpha=0.5,label='Loss curve')
12 plt.legend()


```

```

[ 7.25235662e-05  5.05014458e-01 -1.48671083e-01]

```

Out[76]: &lt;matplotlib.legend.Legend at 0x7f971e77c550&gt;




```

In [77]: 1 from scipy.optimize import curve_fit
2 import matplotlib.pyplot as plt
3
4 def func(k,beta0,beta1,beta2):
5     return (1/(beta0*k+beta1)) + beta2
6 popt, pcov = curve_fit(func, range(1,len(losses2)+1),losses2,p0=[0.001, 0.1,
7 beta0,beta1,beta2 = popt
8 k80.append(popt)
9 print(popt)
10 plt.plot(range(1,len(losses2)+1),[func(k,beta0,beta1,beta2) for k in range(1
11 plt.plot(range(1,len(losses2)+1),losses2,alpha=0.5,label='Loss curve')
12 plt.legend()

```

```
[ 1.03425344e-04  4.81340265e-01 -1.18030695e-01]
```

```
Out[77]: <matplotlib.legend.Legend at 0x7f9722223e90>
```




```

In [78]: 1 from scipy.optimize import curve_fit
2 import matplotlib.pyplot as plt
3
4 def func(k,beta0,beta1,beta2):
5     return (1/(beta0*k+beta1)) + beta2
6 popt, pcov = curve_fit(func, range(1,len(losses3)+1),losses3,p0=[0.001, 0.1,
7 beta0,beta1,beta2 = popt
8 k80.append(popt)
9 print(popt)
10 plt.plot(range(1,len(losses3)+1),[func(k,beta0,beta1,beta2) for k in range(1
11 plt.plot(range(1,len(losses3)+1),losses3,alpha=0.5,label='Loss curve')
12 plt.legend()

```

```
[ 1.46424549e-04  4.71928719e-01 -9.19557722e-02]
```

```
Out[78]: <matplotlib.legend.Legend at 0x7f971e5f5e50>
```




```

In [79]: 1 from scipy.optimize import curve_fit
2 import matplotlib.pyplot as plt
3
4 def func(k,beta0,beta1,beta2):
5     return (1/(beta0*k+beta1)) + beta2
6 popt, pcov = curve_fit(func, range(1,len(losses4)+1),losses4,p0=[0.001, 0.1,
7 beta0,beta1,beta2 = popt
8 k80.append(popt)
9 print(popt)
10 plt.plot(range(1,len(losses4)+1),[func(k,beta0,beta1,beta2) for k in range(1
11 plt.plot(range(1,len(losses4)+1),losses4,alpha=0.5,label='Loss curve')
12 plt.legend()

```

```
[ 1.56356261e-04  4.71935887e-01 -8.90760333e-02]
```


```
Out[79]: <matplotlib.legend.Legend at 0x7f971e5717d0>
```



```
In [80]: 1 from scipy.optimize import curve_fit
2 import matplotlib.pyplot as plt
3
4 def func(k,beta0,beta1,beta2):
5     return (1/(beta0*k+beta1)) + beta2
6 popt, pcov = curve_fit(func, range(1,len(losses5)+1),losses5,p0=[0.001, 0.1,
7 beta0,beta1,beta2 = popt
8 k80.append(popt)
9 print(popt)
10 plt.plot(range(1,len(losses5)+1),[func(k,beta0,beta1,beta2) for k in range(1
11 plt.plot(range(1,len(losses5)+1),losses5,alpha=0.5,label='Loss curve')
12 plt.legend()
```

```
[ 1.55769825e-04  4.69468259e-01 -8.89310787e-02]
```


```
Out[80]: <matplotlib.legend.Legend at 0x7f971e695950>
```





```
In [81]: 1 l_k80 = np.array([[18],[20],[32],[44],[56]])
2 l_p100 = np.array([[18],[20],[32],[44],[56]])
3 l_v100 = np.array([[18],[20],[32],[44],[56]])
```

```
In [82]: 1 b0_k80 = [i[0] for i in k80]
2 b1_k80 = [i[1] for i in k80]
3 b2_k80 = [i[2] for i in k80]
4
5 b0_p100 = [i[0] for i in p100]
6 b1_p100 = [i[1] for i in p100]
7 b2_p100 = [i[2] for i in p100]
8
9 b0_v100 = [i[0] for i in v100]
10 b1_v100 = [i[1] for i in v100]
11 b2_v100 = [i[2] for i in v100]
12
```


```
In [83]: 1 # k80
2 from sklearn.linear_model import LinearRegression as lr
3
4 x = [[i] for i in range(61)]
5
6 kb0 = lr().fit(l_k80,b0_k80)
7 y = kb0.predict(x)
8 plt.plot(x,y)
9 plt.scatter(l_k80,b0_k80)
10 plt.title('beta 0')
11 plt.show()
12
13 kb1 = lr().fit(l_k80,b1_k80)
14 y = kb1.predict(x)
15 plt.plot(x,y)
16 plt.scatter(l_k80,b1_k80)
17 plt.title('beta 1')
18 plt.show()
19
20 kb2 = lr().fit(l_k80,b2_k80)
21 y = kb2.predict(x)
22 plt.plot(x,y)
23 plt.scatter(l_k80,b2_k80)
24 plt.title('beta 2')
25 plt.show()
```







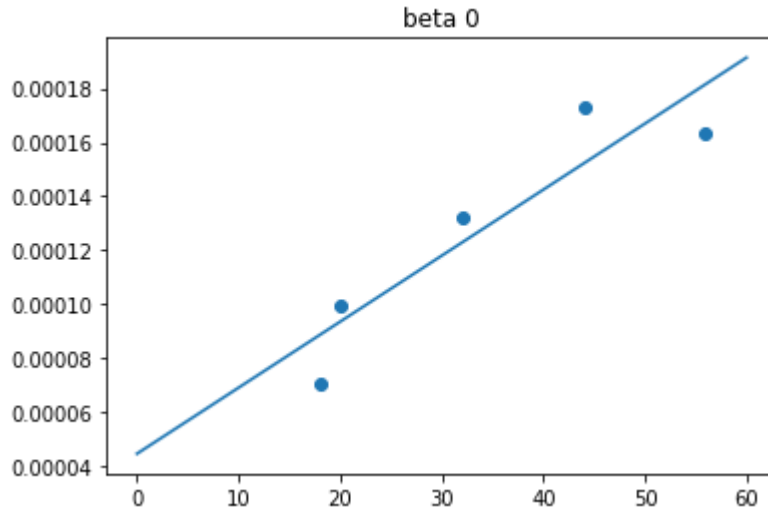



```
In [84]: 1 # p100
2 from sklearn.linear_model import LinearRegression as lr
3
4 x = [[i] for i in range(61)]
5
6 pb0 = lr().fit(l_p100,b0_p100)
7 y = pb0.predict(x)
8 plt.plot(x,y)
9 plt.scatter(l_p100,b0_p100)
10 plt.title('beta 0')
11 plt.show()
12
13 pb1 = lr().fit(l_p100,b1_p100)
14 y = pb1.predict(x)
15 plt.plot(x,y)
16 plt.scatter(l_p100,b1_p100)
17 plt.title('beta 1')
18 plt.show()
19
20 pb2 = lr().fit(l_p100,b2_p100)
21 y = pb2.predict(x)
22 plt.plot(x,y)
23 plt.scatter(l_p100,b2_p100)
24 plt.title('beta 2')
25 plt.show()
```





```
In [85]: 1 # v100
2 from sklearn.linear_model import LinearRegression as lr
3
4 x = [[i] for i in range(61)]
5
6
7 vb0 = lr().fit(l_v100,b0_v100)
8 y = vb0.predict(x)
9 plt.plot(x,y)
10 plt.scatter(l_v100,b0_v100)
11 plt.title('beta 0')
12 plt.show()
13
14 vb1 = lr().fit(l_v100,b1_v100)
15 y = vb1.predict(x)
16 plt.plot(x,y)
17 plt.scatter(l_v100,b1_v100)
18 plt.title('beta 1')
19 plt.show()
20
21 vb2 = lr().fit(l_v100,b2_v100)
22 y = vb2.predict(x)
23 plt.plot(x,y)
24 plt.scatter(l_v100,b2_v100)
25 plt.title('beta 2')
26 plt.show()
```





In [86]:

```

1 # Predictions for ResNet50
2 k80beta = [kb0.predict([[50]]), kb1.predict([[50]]), kb2.predict([[50]])]
3 p100beta = [pb0.predict([[50]]), pb1.predict([[50]]), pb2.predict([[50]])]
4 v100beta = [vb0.predict([[50]]), vb1.predict([[50]]), vb2.predict([[50]])]
5 print(k80beta)
6 print(p100beta)
7 print(v100beta)

```

```

[array([0.00015904]), array([0.46893456]), array([-0.08628932])]
[array([0.00036428]), array([0.45232191]), array([0.0124901])]
[array([0.00016694]), array([0.46555051]), array([-0.08310189])]

```

In [87]:

```

1 losses6 = []
2 textfile = open("loss6.txt", "r")
3 for element in textfile:
4     losses6.append(float(element))
5 textfile.close()


```

```

In [88]: 1 # v100
2
3 l = []
4 err = []
5 epochs = len(losses6)
6 for i in range(epochs):
7     temp = func(i+1,v100beta[0],v100beta[1],v100beta[2])
8     l.append(temp)
9     err.append((temp-losses6[i])/losses6[i])
10 plt.plot([i+1 for i in range(epochs)],losses6,alpha=0.7,color='orange',label=
11 plt.plot([i+1 for i in range(epochs)],l,label='Predicted Loss for V100')
12 plt.legend()
13 # Percentage error
14

```

Out[88]: <matplotlib.legend.Legend at 0x7f9718edc6d0>



```

In [97]: 1 def finv(l,b0,b1,b2):
2         return ((1/(1-b2))-b1)/b0
3         # For V100 Loss at 92% accuracy = 0.27
4         print(finv(0.27,v100beta[0],v100beta[1],v100beta[2]), ' iterations')
5         print('Which is equal to ',finv(0.27,v100beta[0],v100beta[1],v100beta[2]))//3
6

```

[14175.33011273] iterantions  
Which is equal to [36.] epochs

As calculated the model reaches 92% accuracy after epoch 36. Percentage error = 0%

In [88]:

1

In [92]:


1 losses6 = loss\_data['resnet50']

```

In [95]: 1 # p100
2
3 l = []
4 err = []
5 epochs = len(losses6)
6 for i in range(epochs):
7     temp = func((i+1),p100beta[0],p100beta[1],p100beta[2])
8     l.append(temp)
9     err.append((temp-losses6[i])/losses6[i])
10 plt.plot([i+1 for i in range(epochs)],losses6,alpha=0.7,color='orange',label=
11 plt.plot([i+1 for i in range(epochs)],l,label='Predicted Loss for P100')
12 plt.legend()
13 # Percentage error

```

Out[95]: <matplotlib.legend.Legend at 0x7f971dff9450>



```

In [120]: 1 def finv(l,b0,b1,b2):
2           return ((1/(1-b2))-b1)/b0
3 # For P100 Loss at 92% accuracy = 0.18
4 print(finv(0.18,p100beta[0],p100beta[1],p100beta[2]),' iterations')
5 print('Which is equal to ',finv(0.18,p100beta[0],p100beta[1],p100beta[2]))//3
6

```

[15146.13008173] iterations  
Which is equal to [38.] epochs

In [ ]:

1

```

In [90]: 1 losses6 = np.load('resnet50_new.npz')
2 losses6 = losses6.f.a

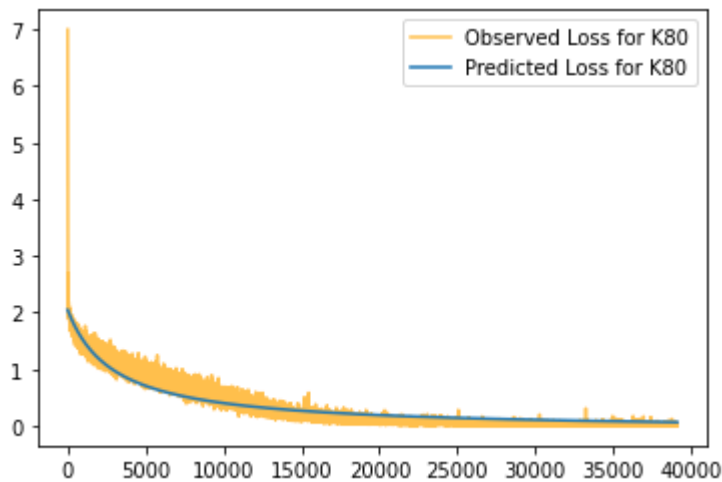
```

```

In [91]: 1 # k80
2
3 l = []
4 err = []
5 epochs = len(losses6)
6 for i in range(epochs):
7     temp = func((i+1),k80beta[0],k80beta[1],k80beta[2])
8     l.append(temp)
9     err.append((temp-losses6[i])/losses6[i])
10 plt.plot([i+1 for i in range(epochs)],losses6[:epochs],alpha=0.7,color='orange')
11 plt.plot([i+1 for i in range(epochs)],l,label='Predicted Loss for K80')
12 plt.legend()
13 # Percentage error

```

Out[91]: <matplotlib.legend.Legend at 0x7f971ee01450>



```

In [118]: 1 def finv(l,b0,b1,b2):
2           return ((1/(1-b2))-b1)/b0
3           # For P100 Loss at 92% accuracy = 0.23
4           print(finv(0.23,k80beta[0],k80beta[1],k80beta[2]),' iterations')
5           print('Which is equal to ',finv(0.23,k80beta[0],k80beta[1],k80beta[2]))//391,
6

```

[16930.7875504] iterations  
Which is equal to [43.] epochs

92% occurs at 36 epochs. error = 19.4%

```

In [ ]: 1

```



```
In [121]: 1 def f(p,w):  
2         return (1.02*(128/w) + 2.78 + 4.92*w/p + 0*w + 0.02*p)**-1
```

```
In [122]: 1 import matplotlib.pyplot as plt  
2  
3 w = range(1,51)  
4 # 92% occurs at epoch 36  
5 epoch = 36  
6 plt.plot(w,[epoch/f(2,w[i]) for i in range(50)],label='2 Parameter Servers')  
7 plt.plot(w,[epoch/f(4,w[i]) for i in range(50)],label='4 Parameter Servers')  
8 plt.xlabel('no of workers')  
9 plt.ylabel('time')  
10 plt.legend()
```

Out[122]: <matplotlib.legend.Legend at 0x7f971e1e3310>

