

# Amrit Parimi

ap4142

## Problem 1

```
In [ ]: 1 !git clone https://github.com/qfgaohtao/pytorch-ssd
```

```
Cloning into 'pytorch-ssd'...
remote: Enumerating objects: 812, done.
remote: Total 812 (delta 0), reused 0 (delta 0), pack-reused 812
Receiving objects: 100% (812/812), 1.05 MiB | 25.60 MiB/s, done.
Resolving deltas: 100% (544/544), done.
```

```
In [ ]: 1 # 1.1
2 !wget -P data http://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar
3 !wget -P data http://pjreddie.com/media/files/VOCtest_06-Nov-2007.tar
4
5 !tar -xvf "data/VOCtrainval_06-Nov-2007.tar" -C "data"
6 !tar -xvf "data/VOCtest_06-Nov-2007.tar" -C "data"
```

```
--2021-12-02 08:03:30-- http://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar (http://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar)
Resolving pjreddie.com (pjreddie.com)... 128.208.4.108
Connecting to pjreddie.com (pjreddie.com)|128.208.4.108|:80... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar (https://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar) [following]
--2021-12-02 08:03:30-- https://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar (https://pjreddie.com/media/files/VOCtrainval_06-Nov-2007.tar)
Connecting to pjreddie.com (pjreddie.com)|128.208.4.108|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 460032000 (439M) [application/octet-stream]
Saving to: 'data/VOCtrainval_06-Nov-2007.tar'

VOCtrainval_06-Nov- 100%[=====] 438.72M 28.8MB/s    in 12s

2021-12-02 08:03:43 (35.2 MB/s) - 'data/VOCtrainval_06-Nov-2007.tar' saved
[460032000/460032000]

URL transformed to HTTPS due to an HSTS header.
```

```
In [ ]: 1 !wget -P models https://storage.googleapis.com/models-hao/mobilenet-v1-ssd-mp-0_675.pth
2 !wget -P models https://storage.googleapis.com/models-hao/voc-model-labels.txt

--2021-12-02 08:04:03-- https://storage.googleapis.com/models-hao/mobilenet-v1-ssd-mp-0_675.pth (https://storage.googleapis.com/models-hao/mobilenet-v1-ssd-mp-0_675.pth)
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.135.128, 74.125.195.128, 74.125.142.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.135.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 37995286 (36M) [application/octet-stream]
Saving to: 'models/mobilenet-v1-ssd-mp-0_675.pth'

mobilenet-v1-ssd-mp 100%[=====] 36.23M 16.9MB/s in 2.1s

2021-12-02 08:04:06 (16.9 MB/s) - 'models/mobilenet-v1-ssd-mp-0_675.pth' saved [37995286/37995286]

--2021-12-02 08:04:08-- https://storage.googleapis.com/models-hao/voc-model-labels.txt (https://storage.googleapis.com/models-hao/voc-model-labels.txt)
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.135.128, 74.125.195.128, 74.125.142.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.135.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 145 [text/plain]
Saving to: 'models/voc-model-labels.txt'

voc-model-labels.txt 100%[=====] 145 --.-KB/s in 0s

2021-12-02 08:04:08 (210 MB/s) - 'models/voc-model-labels.txt' saved [145/145]
```

```
In [ ]: 1 !python pytorch-ssd/eval_ssd.py --net mb1-ssd \
2                                     --trained_model models/mobilenet_v2-1.0_20190128-073700
3                                     --dataset_type voc \
4                                     --dataset /content/data/VOCdevkit \
5                                     --label_file /content/models/voc_label_map.pbtxt
boat: 0.5612220055063379
bottle: 0.34852207089954895
bus: 0.7677814849265677
car: 0.7280986468467315
cat: 0.8369208203985581
chair: 0.5169138632991064
cow: 0.6238697603075337
diningtable: 0.7062172972736019
dog: 0.7872868219961326
horse: 0.819446325939355
motorbike: 0.7918539457195842
person: 0.7023619139901851
pottedplant: 0.39852951468542563
sheep: 0.6066678298227772
sofa: 0.7573083661544429
train: 0.8262441264750008
tvmonitor: 0.6461898726506375

Average Precision Across All Classes:0.6759028864961053
```

```
In [ ]: 1
```

```
In [ ]: 1
```

In [ ]:

```
1 # 1.2
2 # !wget -P models https://storage.googleapis.com/models-hao/gun_model_2.21.pth
3 !wget -P models https://storage.googleapis.com/models-hao/open-images-model-1-labels.txt

--2021-12-02 09:38:23-- https://storage.googleapis.com/models-hao/gun_model_2.21.pth (https://storage.googleapis.com/models-hao/gun_model_2.21.pth)
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.135.128, 74.125.195.128, 74.125.142.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.135.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 27044080 (26M) [application/octet-stream]
Saving to: 'models/gun_model_2.21.pth.1'

gun_model_2.21.pth. 100%[=====] 25.79M 170MB/s in 0.2s

2021-12-02 09:38:24 (170 MB/s) - 'models/gun_model_2.21.pth.1' saved [27044080/27044080]

--2021-12-02 09:38:24-- https://storage.googleapis.com/models-hao/open-images-model-labels.txt (https://storage.googleapis.com/models-hao/open-images-model-labels.txt)
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.135.128, 74.125.195.128, 74.125.142.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.135.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 26 [text/plain]
Saving to: 'models/open-images-model-labels.txt.1'

open-images-model-1 100%[=====] 26 ---KB/s in 0s

2021-12-02 09:38:24 (30.0 MB/s) - 'models/open-images-model-labels.txt.1' saved [26/26]
```

```
In [ ]: 1 !pip install boto3
```

```
Collecting boto3
  Downloading boto3-1.20.18-py3-none-any.whl (131 kB)
    |██████████| 131 kB 8.6 MB/s
Collecting botocore<1.24.0,>=1.23.18
  Downloading botocore-1.23.18-py3-none-any.whl (8.4 MB)
    |██████████| 8.4 MB 33.1 MB/s
Collecting s3transfer<0.6.0,>=0.5.0
  Downloading s3transfer-0.5.0-py3-none-any.whl (79 kB)
    |██████████| 79 kB 10.6 MB/s
Collecting jmespath<1.0.0,>=0.7.1
  Downloading jmespath-0.10.0-py2.py3-none-any.whl (24 kB)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in /opt/conda/lib/python3.7/site-packages (from botocore<1.24.0,>=1.23.18->boto3) (2.8.2)
Requirement already satisfied: urllib3<1.27,>=1.25.4 in /opt/conda/lib/python3.7/site-packages (from botocore<1.24.0,>=1.23.18->boto3) (1.26.7)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.7/site-packages (from python-dateutil<3.0.0,>=2.1->botocore<1.24.0,>=1.23.18->boto3) (1.16.0)
Installing collected packages: jmespath, botocore, s3transfer, boto3
Successfully installed boto3-1.20.18 botocore-1.23.18 jmespath-0.10.0 s3transfer-0.5.0
```

```
In [ ]: 1 !python pytorch-ssd/open_images_downloader.py --root ~/data/open_images --cl
```

```
2021-12-02 10:57:58,075 - root - Download https://storage.googleapis.com/openimages/2018\_04/class-descriptions-boxable.csv. (https://storage.googleapis.com/openimages/2018\_04/class-descriptions-boxable.csv)
2021-12-02 10:57:58,119 - root - Download https://storage.googleapis.com/openimages/2018\_04/train/train-annotations-bbox.csv. (https://storage.googleapis.com/openimages/2018\_04/train/train-annotations-bbox.csv)
2021-12-02 10:58:08,626 - root - Read annotation file /home/parim/data/open_images/train-annotations-bbox.csv
```

```
In [ ]: 1 !python pytorch-ssd/train_ssd.py --dataset_type open_images --datasets ~/dat
          r-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is
          deprecated. If you meant to do this, you must specify 'dtype=object' when c
          reating the ndarray
          mode = random.choice(self.sample_options)
/content/pytorch-ssd/vision/transforms/transforms.py:247: VisibleDeprecation
Warning: Creating an ndarray from ragged nested sequences (which is a list-o
r-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is
          deprecated. If you meant to do this, you must specify 'dtype=object' when c
          reating the ndarray
          mode = random.choice(self.sample_options)
/content/pytorch-ssd/vision/transforms/transforms.py:247: VisibleDeprecation
Warning: Creating an ndarray from ragged nested sequences (which is a list-o
r-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is
          deprecated. If you meant to do this, you must specify 'dtype=object' when c
          reating the ndarray
          mode = random.choice(self.sample_options)
2021-12-01 20:54:53,058 - root - INFO - Epoch: 88, Step: 100, Average Loss:
          1.8540, Average Regression Loss 0.5043, Average Classification Loss: 1.3497
/content/pytorch-ssd/vision/transforms/transforms.py:247: VisibleDeprecation
Warning: Creating an ndarray from ragged nested sequences (which is a list-o
```

```
In [ ]: 1 # Before finetuning
2 !python pytorch-ssd/eval_ssd.py --net mb1-ssd \
          --trained_model models/mb1-ssd
          --dataset_type open_images \
          --dataset /root/data/open_imag
          --label_file /content/models/o
          .uction... 0.040577 seconds.
process image 120
Load Image: 0.030503 seconds.
Inference time: 0.010792016983032227
Prediction: 0.044940 seconds.
process image 121
Load Image: 0.035109 seconds.
Inference time: 0.011407136917114258
Prediction: 0.043935 seconds.
process image 122
Load Image: 0.020850 seconds.
Inference time: 0.011580944061279297
Prediction: 0.049960 seconds.

Average Precision Per-class:
Handgun: 0.7526829425387441
Shotgun: 0.3529092048726733

Average Precision Across All Classes: 0.5527960737057087
```

```
In [ ]: 1 # After finetuning
2 !python pytorch-ssd/eval_ssd.py --net mb1-ssd \
3                                         --trained_model models/mb1-ssd
4                                         --dataset_type open_images \
5                                         --dataset /root/data/open_imag
6                                         --label_file /content/models/o
Prediction: 0.029038 seconds.
process image 120
Load Image: 0.036938 seconds.
Inference time: 0.01143193244934082
Prediction: 0.041948 seconds.
process image 121
Load Image: 0.031048 seconds.
Inference time: 0.01518869400024414
Prediction: 0.038679 seconds.
process image 122
Load Image: 0.017745 seconds.
Inference time: 0.011113643646240234
Prediction: 0.029559 seconds.

Average Precision Per-class:
Handgun: 0.8130237277657372
Shotgun: 0.5787321045226184

Average Precision Across All Classes: 0.6958779161441777
```

```
In [ ]: 1
```

```
In [1]: 1 !pip install onnx
```

```
Collecting onnx
  Downloading onnx-1.10.2-cp37-cp37m-manylinux_2_12_x86_64.manylinux2010_x86_6
  4.whl (12.7 MB)
    |██████████| 12.7 MB 6.7 MB/s
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (f
rom onnx) (1.15.0)
Requirement already satisfied: typing-extensions>=3.6.2.1 in /usr/local/lib/pyt
hon3.7/dist-packages (from onnx) (3.10.0.2)
Requirement already satisfied: protobuf in /usr/local/lib/python3.7/dist-pacak
es (from onnx) (3.17.3)
Requirement already satisfied: numpy>=1.16.6 in /usr/local/lib/python3.7/dist-p
ackages (from onnx) (1.19.5)
Installing collected packages: onnx
Successfully installed onnx-1.10.2
```

```
In [ ]: 1 # Converting to ONNX
2 # 1.3
3 !python pytorch-ssd/convert_to_caffe2_models.py mb1-ssd /content/models/mb1-
[▶]
```

/usr/local/lib/python3.7/dist-packages/caffe2/python/onnx/backend.py:878: UserWarning: ShapeInferenceWarning: utils module not found in ONNX version 1.10.2  
    warnings.warn("ShapeInferenceWarning: utils module not found in ONNX version  
    {}".format(onnx.\_\_version\_\_))  
/usr/local/lib/python3.7/dist-packages/caffe2/python/onnx/backend.py:885: UserWarning: OptimizerWarning: optimizer module not found in ONNX version 1.10.2  
    warnings.warn("OptimizerWarning: optimizer module not found in ONNX version  
    {}".format(onnx.\_\_version\_\_))  
Save the model in binary format to the files models/mb1-ssd\_init\_net.pb and models/mb1-ssd\_predict\_net.pb.  
Save the model in txt format to the files models/mb1-ssd\_init\_net.pbtxt and models/mb1-ssd\_predict\_net.pbtxt.

```
In [2]: 1 !git clone https://github.com/onnx/onnx.git
```

Cloning into 'onnx'...  
remote: Enumerating objects: 29055, done.  
remote: Counting objects: 100% (868/868), done.  
remote: Compressing objects: 100% (492/492), done.  
remote: Total 29055 (delta 504), reused 605 (delta 361), pack-reused 28187  
Receiving objects: 100% (29055/29055), 20.78 MiB | 21.15 MiB/s, done.  
Resolving deltas: 100% (16131/16131), done.

```
In [ ]: 1 !git clone https://gitlab.com/graphviz/graphviz.git
```

Cloning into 'graphviz'...  
remote: Enumerating objects: 106337, done.  
remote: Counting objects: 100% (1818/1818), done.  
remote: Compressing objects: 100% (108/108), done.  
remote: Total 106337 (delta 1744), reused 1753 (delta 1710), pack-reused 104519  
Receiving objects: 100% (106337/106337), 176.20 MiB | 25.44 MiB/s, done.  
Resolving deltas: 100% (80825/80825), done.

```
In [3]: 1 !pip install onnx  
2 !pip install pydot
```

```
Requirement already satisfied: onnx in /usr/local/lib/python3.7/dist-packages  
(1.10.2)  
Requirement already satisfied: typing-extensions>=3.6.2.1 in /usr/local/lib/python3.7/dist-packages (from onnx) (3.10.0.2)  
Requirement already satisfied: protobuf in /usr/local/lib/python3.7/dist-packages (from onnx) (3.17.3)  
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from onnx) (1.15.0)  
Requirement already satisfied: numpy>=1.16.6 in /usr/local/lib/python3.7/dist-packages (from onnx) (1.19.5)  
Requirement already satisfied: pydot in /usr/local/lib/python3.7/dist-packages (1.3.0)  
Requirement already satisfied: pyparsing>=2.1.4 in /usr/local/lib/python3.7/dist-packages (from pydot) (3.0.6)
```

```
In [ ]: 1 !git clone https://github.com/onnx/tutorials.git
```

```
Cloning into 'tutorials'...  
remote: Enumerating objects: 654, done.  
remote: Counting objects: 100% (65/65), done.  
remote: Compressing objects: 100% (57/57), done.  
remote: Total 654 (delta 17), reused 17 (delta 5), pack-reused 589  
Receiving objects: 100% (654/654), 37.26 MiB | 32.09 MiB/s, done.  
Resolving deltas: 100% (289/289), done.
```

```
In [6]: 1 import os  
2 !python onnx/onnx/tools/net_drawer.py --input "mb1-ssd.onnx" --output squeez
```

```
In [7]: 1 !dot -Tsvg squeezenet.dot -o squeezenet.svg
```

```
In [ ]: 1 !pip install pillow
```

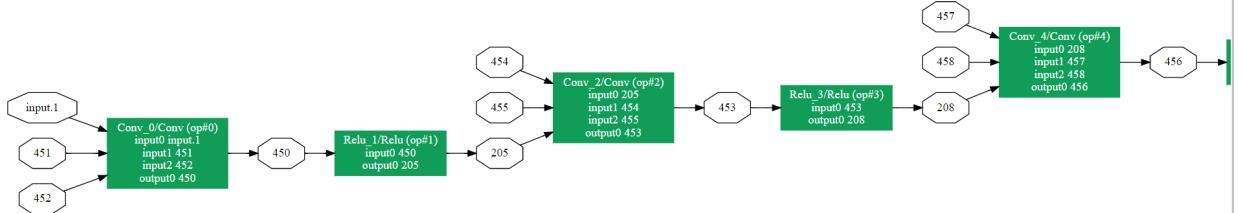
```
Requirement already satisfied: pillow in /opt/conda/lib/python3.7/site-packages  
(8.4.0)
```

```
In [8]: 1 from IPython.display import SVG, display  
2 display(SVG('squeezenet.svg'))
```

```
<IPython.core.display.SVG object>
```

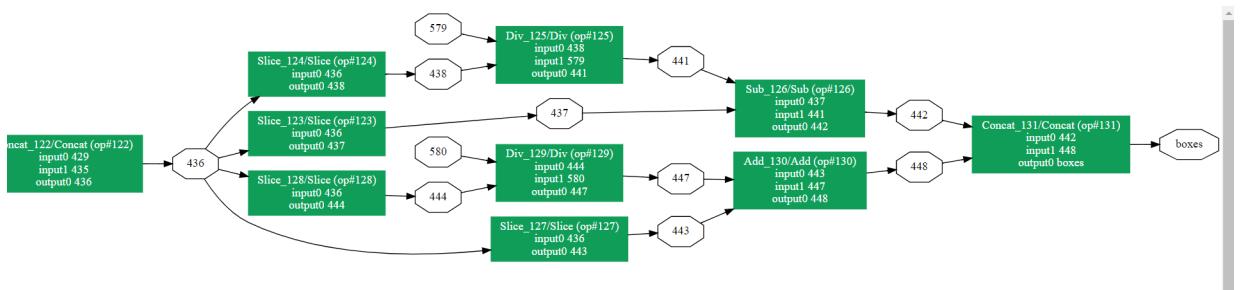
```
In [2]: 1 # As I cannot view the output of net_drawer in Jupyter notebook I have included
          from PIL import Image
          Image.open('C:/Users/parim/Desktop/Columbia Class/DL/hw5/net-start.PNG')
```

Out[2]:



```
In [3]: 1 Image.open('C:/Users/parim/Desktop/Columbia Class/DL/hw5/net-end.PNG')
```

Out[3]:



```
In [ ]: 1
```

```
In [ ]: 1 # 1.5,1.6
          2 !pip install onnxruntime
```

```
In [ ]: 1
          2 !sudo docker pull mcr.microsoft.com/onnxruntime/server
```

Using default tag: latest  
latest: Pulling from onnxruntime/server  
Digest: sha256:a23da0977bbc4aca4d3de56ad648ebde86031e61d7a3b7cbe1daaebbc7f6ad3d  
Status: Image is up to date for mcr.microsoft.com/onnxruntime/server:latest  
mcr.microsoft.com/onnxruntime/server:latest

```
In [ ]: 1 !sudo docker run -it -v $(pwd):$(pwd) -p 9005:8001 mcr.microsoft.com/onnxrun
[2021-12-02 11:59:52.246] [ServerApp] [info] [ onnxruntime session_state_initializer.cc:345 SaveInputOutputNamesToNodeMapping]: Graph input with name backbone.model.layer2.0.4.bn2.num_batches_tracked is not used by any node.
[2021-12-02 11:59:52.246] [ServerApp] [info] [ onnxruntime session_state_initializer.cc:345 SaveInputOutputNamesToNodeMapping]: Graph input with name backbone.model.layer2.0.5.bn1.num_batches_tracked is not used by any node.
[2021-12-02 11:59:52.246] [ServerApp] [info] [ onnxruntime session_state_initializer.cc:345 SaveInputOutputNamesToNodeMapping]: Graph input with name backbone.model.layer2.0.5.bn2.num_batches_tracked is not used by any node.
[2021-12-02 11:59:52.246] [ServerApp] [info] [ onnxruntime inference_session.cc:620 Initialize]: Session successfully initialized.
[2021-12-02 11:59:52.248] [ServerApp] [info] GRPC Listening at: 0.0.0.0:5005
1
[2021-12-02 11:59:52.248] [ServerApp] [info] Listening at: http://0.0.0.0:8001 (http://0.0.0.0:8001)
[2021-12-02 12:08:33.415] [ac257471-add5-47a6-9d53-7ff6d99a95c2] [info] Mode 1 Name: default, Version: 1, Action: predict
[2021-12-02 12:09:05.223] [cd2768d4-f409-438a-b2cf-a0bca1199302] [info] Mode 1 Name: default, Version: 1, Action: predict
^C
```

```
In [ ]: 1 import os
2
3 import tutorials.tutorials.assets.onnx_ml_pb2 as onnx_ml_pb2
4 import tutorials.tutorials.assets.predict_pb2 as predict_pb2
```

```
In [ ]: 1 import tensorflow as tf
2 x = tf.cast(img2, tf.float32)
3 img2 = tf.keras.applications.mobilenet.preprocess_input(x)
```

```
In [ ]: 1 !pip install tensorflow
```

...

In [ ]:

```
1 import numpy as np
2 import requests
3 from PIL import Image
4 import matplotlib.pyplot as plt
5 import matplotlib.patches as patches
6 img1 = Image.open("ims/831f04c7b16ad888.jpg")
7 img1 = img1.resize((300, 300), Image.BILINEAR)
8 img1
9
```

Out[39]:



In [ ]:

```
1 img_data = np.array(img1)
2 img_data = np.transpose(img_data, [2, 0, 1])
3 img_data = np.expand_dims(img_data, 0)
4 mean_vec = np.array([0.5093532301981356, 0.47697768609685176, 0.443410724237]
5 stddev_vec = np.array([0.17204669910914244, 0.1762167087185571, 0.1910801545]
6 norm_img_data = np.zeros(img_data.shape).astype('float32')
7 for i in range(img_data.shape[1]):
8     norm_img_data[:,i,:,:] = (img_data[:,i,:,:]/255 - mean_vec[i]) / stddev_
```

```
In [ ]: 1 input_tensor = onnx_ml_pb2.TensorProto()
2 input_tensor.dims.extend(norm_img_data.shape)
3 input_tensor.data_type = 1
4 input_tensor.raw_data = norm_img_data.tobytes()
5
6 request_message = predict_pb2.PredictRequest()
7
8 print(request_message.inputs)
9 # For your model, the inputs name should be something else customized by you
10 request_message.inputs["input.1"].data_type = input_tensor.data_type
11 request_message.inputs["input.1"].dims.extend(input_tensor.dims)
12 request_message.inputs["input.1"].raw_data = input_tensor.raw_data
13
14 content_type_headers = ['application/x-protobuf', 'application/octet-stream']
15
16 for h in content_type_headers:
17     request_headers = {
18         'Content-Type': h,
19         'Accept': 'application/x-protobuf'
20     }
21
22 }
```

```
In [ ]: 1 PORT_NUMBER = 9005 # Change appropriately if needed based on any changes whe  
2 inference_url = "http://127.0.0.1:" + str(PORT_NUMBER) + "/v1/models/default  
3 response = requests.post(inference_url, headers=request_headers, data=reques
```

```
In [ ]: 1 # print(response.content)
2 response_message = predict_pb2.PredictResponse()
3 response_message.ParseFromString(response.content)
4
5 # print(response_message.outputs)
6 # For your model, the outputs names should be something else customized by y
7 bboxes = np.frombuffer(response_message.outputs['boxes'].raw_data, dtype=np.
8 # labels = np.frombuffer(response_message.outputs['labels'].raw_data, dtype=
9 scores = np.frombuffer(response_message.outputs['scores'].raw_data, dtype=np.
10
11 print(scores)
12 print('Boxes shape:', response_message.outputs['boxes'].dims)
13 # print('Labels shape:', response_message.outputs['labels'].dims)
14 print('Scores shape:', response_message.outputs['scores'].dims)
```

```
[0.9897613  0.00536984 0.00486883 ... 0.7378382  0.25728315 0.00487863]  
Boxes shape: [1, 3000, 4]  
Scores shape: [1, 3000, 3]
```

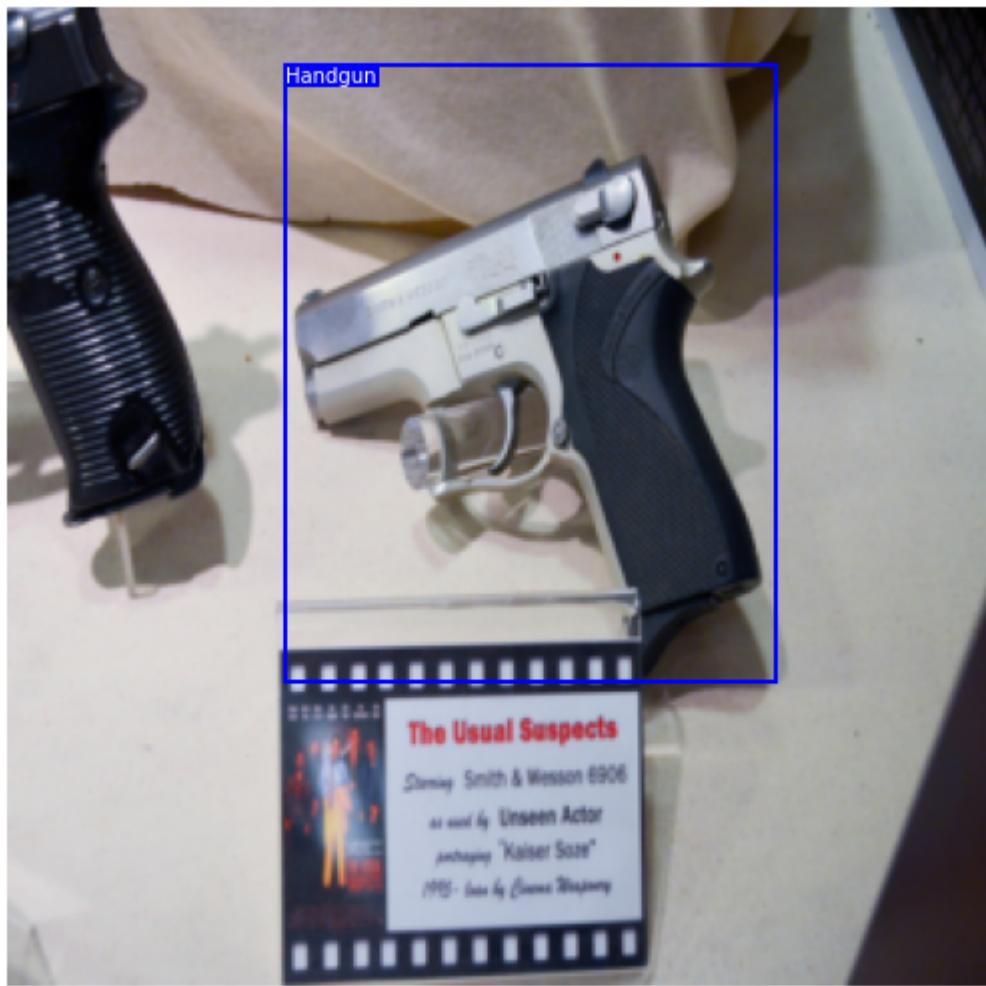
In [ ]:

```
1 print(bboxes[0],bboxes[1],bboxes[2],bboxes[3])
2 print(len(bboxes),len(scores))
3 print(sum(scores>0.999),max(scores))
4 b = []
5 for i in range(len(scores)//3):
6     if max([scores[3*i+1]])>0.99:
7         b.append([bboxes[4*i],bboxes[4*i+1],bboxes[4*i+2],bboxes[4*i+3]])
8         print([scores[3*i],scores[3*i+1],scores[3*i+2]])
9 print(b)
10 print(len(b))
11 print([[300*i for i in j] for j in b])
```

```
-0.020249806 -0.011044905 0.10776336 0.107537776
12000 9000
600 0.99998856
[0.008273488, 0.99049395, 0.0012325967]
[[0.2823299, 0.059160978, 0.7833177, 0.68726707]]
1
[[84.69896614551544, 17.7482932806015, 234.9953055381775, 206.18011951446533]]
```

```
In [ ]: 1 # Parse the list of class labels
2 classes = [line.rstrip('\n') for line in open('models/open-images-model-labe
3 # classes = ['Handgun', 'Shotgun']
4
5 # Plot the bounding boxes on the image
6 plt.figure()
7 fig, ax = plt.subplots(1, figsize=(12,9))
8 ax.imshow(img1)
9
10 resized_width = 300 # we resized the original image, remember ?
11 resized_height = 300
12 num_boxes = len(b) # we limit displaying to just 10 boxes to avoid clogging
13 # The results are already sorted based on box confidences, so
14
15 for c in range(num_boxes):
16     y1, x1, y2, x2 = b[c][0] * resized_height, b[c][1] * resized_width, b[c]
17     color = 'blue'
18     box_h = (y2 - y1)
19     box_w = (x2 - x1)
20     bbox = patches.Rectangle((y1, x1), box_h, box_w, linewidth=2, edgecolor=
21     ax.add_patch(bbox)
22     # print(labels[c] - 1)
23     plt.text(y1, x1, s=classes[1], color='white', verticalalignment='top', b
24     plt.axis('off')
25
26 # Save image
27 # plt.savefig("output/ssd_result.jpg", bbox_inches='tight', pad_inches=0.0)
28 plt.show()
```

<Figure size 432x288 with 0 Axes>



In [ ]:

1

In [ ]:

1

```
In [ ]: 1 img2 = Image.open("ims/bf7b28e33c453f4a.jpg")
2 img2 = img2.resize((300, 300), Image.BILINEAR)
3 img2
```

Out[64]:



```
In [ ]: 1 img_data = np.array(img2)
2 img_data = np.transpose(img_data, [2, 0, 1])
3 img_data = np.expand_dims(img_data, 0)
4 mean_vec = np.array([0.5093532301981356, 0.47697768609685176, 0.443410724237]
5 stddev_vec = np.array([0.17204669910914244, 0.1762167087185571, 0.1910801545]
6 norm_img_data = np.zeros(img_data.shape).astype('float32')
7 for i in range(img_data.shape[1]):
8     norm_img_data[:,i,:,:] = (img_data[:,i,:,:]/255 - mean_vec[i]) / stddev_
```

```
In [ ]: 1 input_tensor = onnx_ml_pb2.TensorProto()
2 input_tensor.dims.extend(norm_img_data.shape)
3 input_tensor.data_type = 1
4 input_tensor.raw_data = norm_img_data.tobytes()
5
6 request_message = predict_pb2.PredictRequest()
7
8 print(request_message.inputs)
9 # For your model, the inputs name should be something else customized by you
10 request_message.inputs["input.1"].data_type = input_tensor.data_type
11 request_message.inputs["input.1"].dims.extend(input_tensor.dims)
12 request_message.inputs["input.1"].raw_data = input_tensor.raw_data
13
14 content_type_headers = ['application/x-protobuf', 'application/octet-stream']
15
16 for h in content_type_headers:
17     request_headers = {
18         'Content-Type': h,
19         'Accept': 'application/x-protobuf'
20     }
```

{}

```
In [ ]: 1 PORT_NUMBER = 9005 # Change appropriately if needed based on any changes when
2 inference_url = "http://127.0.0.1:" + str(PORT_NUMBER) + "/v1/models/default"
3 response = requests.post(inference_url, headers=request_headers, data=reques
```

```
In [ ]: 1 # print(response.content)
2 response_message = predict_pb2.PredictResponse()
3 response_message.ParseFromString(response.content)
4
5 # print(response_message.outputs)
6 # For your model, the outputs names should be something else customized by you
7 bboxes = np.frombuffer(response_message.outputs['boxes'].raw_data, dtype=np.
8 # Labels = np.frombuffer(response_message.outputs['Labels'].raw_data, dtype=
9 scores = np.frombuffer(response_message.outputs['scores'].raw_data, dtype=np.
10
11 print(scores)
12 print('Boxes shape:', response_message.outputs['boxes'].dims)
13 # print('Labels shape:', response_message.outputs['Labels'].dims)
14 print('Scores shape:', response_message.outputs['scores'].dims)
```

[0.9965726 0.00137203 0.00205533 ... 0.9824255 0.00197902 0.01559554]  
 Boxes shape: [1, 3000, 4]  
 Scores shape: [1, 3000, 3]

```
In [ ]: 1 print(bboxes[0],bboxes[1],bboxes[2],bboxes[3])
2 print(len(bboxes),len(scores))
3 print(sum(scores>0.999),max(scores))
4 b = []
5 for i in range(len(scores)//3):
6     if max([scores[3*i+2]])>0.96:
7         b.append([bboxes[4*i],bboxes[4*i+1],bboxes[4*i+2],bboxes[4*i+3]])
8         print([scores[3*i],scores[3*i+1],scores[3*i+2]])
9 print(b)
10 print(len(b))
11 print([[300*i for i in j] for j in b])
```

-0.020073365 -0.010506172 0.1116215 0.09724927  
 12000 9000  
 910 1.0  
 [0.024588618, 0.008374293, 0.96703714]  
 [[0.14970711, 0.27332062, 0.7277031, 0.707513]]  
 1  
 [[44.91213262081146, 81.99618458747864, 218.31092834472656, 212.2538924217224]]

In [ ]:

```
1 # Parse the list of class labels
2 classes = [line.rstrip('\n') for line in open('models/open-images-model-labe
3 # classes = ['Handgun', 'Shotgun']
4
5 # Plot the bounding boxes on the image
6 plt.figure()
7 fig, ax = plt.subplots(1, figsize=(12,9))
8 ax.imshow(img2)
9
10 resized_width = 300 # we resized the original image, remember ?
11 resized_height = 300
12 num_boxes = len(b) # we limit displaying to just 10 boxes to avoid clogging
13 # The results are already sorted based on box confidences, so
14
15 for c in range(num_boxes):
16     y1, x1, y2, x2 = b[c][0] * resized_height, b[c][1] * resized_width, b[c]
17     color = 'blue'
18     box_h = (y2 - y1)
19     box_w = (x2 - x1)
20     bbox = patches.Rectangle((y1, x1), box_h, box_w, linewidth=2, edgecolor=
21     ax.add_patch(bbox)
22 #     print(labels[c] - 1)
23     plt.text(y1, x1, s=classes[2], color='white', verticalalignment='top', b
24     plt.axis('off')
25
26 # Save image
27 # plt.savefig("output/ssd_result.jpg", bbox_inches='tight', pad_inches=0.0)
28 plt.show()
```

&lt;Figure size 432x288 with 0 Axes&gt;



In [ ]:

1

In [ ]:

1

## Problem 2

```
1 1. DL Frameworks supported and their versions:  
2  
3 IBM, Google, Microsoft, and Amazon  
4  
5 a) IBM :- IBM Cloud Paks 4.0:<br>  
6 i. Pytorch version - 1.7<br>  
7 ii. Tensorflow version - 2.4<br>  
8 Other useful ML frameworks: XGBoost version(1.3), Scikit-learn(0.23),  
Spark version(3.0)  
9  
10 b) Google :- GCP:<br>  
11 i. Tensorflow version - upto 2.7 (Long term vevrsion support for 2.6)<br>  
12 ii. Pytorch version - 1.10<br>  
13 iii. MXNet version - 1.5.1<br>  
14 iv. XGBoost version - 1.4.2<br>  
15 v. Caffe version - 1.0<br>  
16 Other frameworks include - Pytorch-XLA, RAPIDS, Chainer, CNTK, etc.
```

```
17
18 c) Microsoft :- Azure:<br>
19 i. Pytorch version - 1.9.0<br>
20 ii. Tensorflow version - 2.5<br>
21 iii. Chainer version - 7.0.0<br>
22 Other frameworks include - Horovod(0.21.3), CUDA/cudNN(11), CNTK, scikit-
23 learn
24
25 d) Amazon :- AWS:<br>
26 i. Pytorch version - 1.6<br>
27 ii. Tensorflow version - 2.1<br>
28 iii. Keras version - 2.2.4<br>
29 iv. ONNX version - 1.7.0<br>
30 v. MXNet version - 1.8.0<br>
31 vi. Caffe version - 1.0<br>
32 Other frameworks include - XGBoost(1.3.3), Gluon toolkit, CNTK(2.4),
33 Theano(1.0)
34
35 2. Compute Units (GPUs):
36
37 a) IBM :- Nvidia K80, V100, P100, T4<br>
38 b) Google :- Nvidia K80, P4, V100, A100, T4, P100<br>
39 c) Microsoft :- Nvidia V100, P40, M60, K80, T4, P100<br>
40 d) Amazon :- Nvidia A100, V100, K80, T4, M60, A10G, AMD Radeon Pro V520
41
42 3. Model lifecycle management:
43
44 a) IBM :- IBM Cloud Pak<br>
45 b) Google :- Google cloud interface with KubeFlow<br>
46 c) Microsoft :- MLOps<br>
47 d) Amazon :- Amazon SageMaker
48
49 4. Monitoring:
50
51 a) IBM :- IBM Cloud Application Performance Management (APM), LogDNA,
52 Sysdig<br>
53 b) Google :- Google Cloud's operations suite<br>
54 c) Microsoft :- Azure Monitor<br>
55 d) Amazon :- Amazon CloudWatch
56
57 5. Visualization:
58
59 a) IBM :- Watson OpenScale<br>
60 b) Google :- Tensorboard while using Tensorflow<br>
61 c) Microsoft :- Azure Monitor Metrics<br>
62 d) Amazon :- Amazon SageMaker has built-in support for metrics
63
64 6. Elastic Scaling:
65
66 a) IBM :- AutoScale can be used for elastic scaling<br>
67 b) Google :- Google Cloud offers server-side load balancing and
68 autoscaling for groups of instances so we can distribute incoming traffic
69 across multiple virtual machine instances<br>
```

```
65 c) Microsoft :- Azure virtual machine scale set can automatically increase  
or decrease the number of VM instances that run your application.  
Autoscale is a built-in feature of Cloud Services, Mobile Services,  
Virtual Machines, and Websites that helps applications perform their best  
when demand changes. For example, you could have a web app that handles  
millions of requests during the day and none at night. Autoscale can scale  
your service by any of these—or by a custom metric you define.<br>  
66 d) Amazon :- AWS Auto Scaling lets you build scaling plans that automate  
how groups of different resources respond to changes in demand. You can  
optimize availability, costs, or a balance of both.  
67  
68 7. Training job description:  
69  
70 a) IBM :- <br>  
71 model_definition:  
72   framework:  
73     name: tensorflow  
74     version: "1.15"  
75   runtimes:  
76     name: python  
77     version: "3.5"  
78   name: tf-mnist  
79   description: Simple MNIST model implemented in TF  
80   execution:  
81     command: python3 convolutional_network.py --trainImagesFile  
${DATA_DIR}/train-images-idx3-ubyte.gz  
     --trainLabelsFile ${DATA_DIR}/train-labels-idx1-ubyte.gz --  
testImagesFile ${DATA_DIR}/t10k-images-idx3-ubyte.gz  
     --testLabelsFile ${DATA_DIR}/t10k-labels-idx1-ubyte.gz --  
learningRate 0.001  
     --trainingIters 2000000  
   compute_configuration:  
85     name: v100  
86 training_data_reference:  
87   name: MNIST image data files  
88   connection:  
89     endpoint_url: <auth-url>  
90     access_key_id: <username>  
91     secret_access_key: <password>  
92   source:  
93     bucket: mnist-training-data  
94   type: s3  
95 training_results_reference:  
96   name: DL Model Storage  
97   connection:  
98     endpoint_url: <auth-url>  
99     access_key_id: <username>  
100    secret_access_key: <password>  
101 target:  
102   bucket: mnist-training-models  
103   type: s3  
104  
105  
106  
107  
108  
109 b) Google :- <br>  
110 PACKAGE_PATH="/path/to/your/application/sources"
```

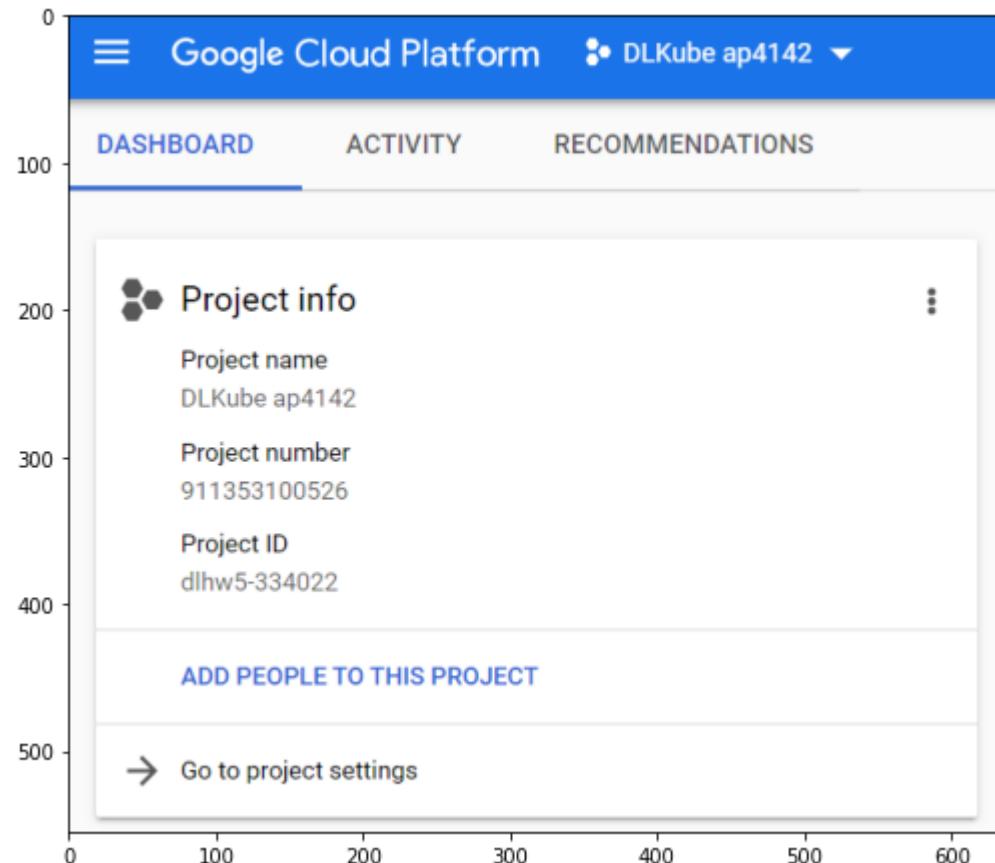
```
111 now=(date +"%Y%m%d_%H%M%S")
112 JOB_NAME="your_name_now"
113 MODULE_NAME="trainer.task"
114 JOB_DIR="gs://your/chosen/job/output/path"
115 REGION="us-east1"
116 RUNTIME_VERSION="2.6"
117
118 gcloud ai-platform jobs submit training JOB_NAME \
119     --scale-tier basic \
120     --package-path PACKAGE_PATH \
121     --module-name MODULE_NAME \
122     --job-dir JOB_DIR \
123     --region REGION \
124     -- \
125     --user_first_arg=first_arg_value \
126     --user_second_arg=second_arg_value
127
128
129
130
131 c) Microsoft :- <br>
132 from azureml.core import ScriptRunConfig, Experiment
133
134     experiment = Experiment(workspace, 'MyExperiment')
135     cluster = workspace.compute_targets['MyCluster']
136     env = Environment.get(ws, name='MyEnvironment')
137     config = ScriptRunConfig(source_directory='.',
138                             script='train.py',
139                             arguments=['--arg1', arg1_val, '--arg2',
140                                         arg2_val],
141                                         compute_target=cluster,
142                                         environment=env)
143     script_run = experiment.submit(config)
144
145
146 d) Amazon :-
147     create-training-job
148     --training-job-name <value>
149     [--hyper-parameters <value>]
150     --algorithm-specification <value>
151     --role-arn <value>
152     [--input-data-config <value>]
153     --output-data-config <value>
154     --resource-config <value>
155     [--vpc-config <value>]
156     --stopping-condition <value>
157     [--tags <value>]
158     [--enable-network-isolation | --no-enable-network-isolation]
159     [--enable-inter-container-traffic-encryption | --no-enable-inter-
160     container-traffic-encryption]
161     [--enable-managed-spot-training | --no-enable-managed-spot-training]
162     [--checkpoint-config <value>]
163     [--debug-hook-config <value>]
164     [--debug-rule-configurations <value>]
165     [--tensor-board-output-config <value>]
166     [--experiment-config <value>]
```

```
166 [--profiler-config <value>]
167 [--profiler-rule-configurations <value>]
168 [--environment <value>]
169 [--retry-strategy <value>]
170 [--cli-input-json <value>]
171 [--generate-cli-skeleton <value>]
172
173
174 As we can see the different cloud platforms require similar data like the
framaework to be used, the pythhon file to be run, hyperparameter values,
locatioin of thhe file, environment or workspace to be used, package path,
stopoping condition or number of iterations or time to be run.
175
```

## Problem 3

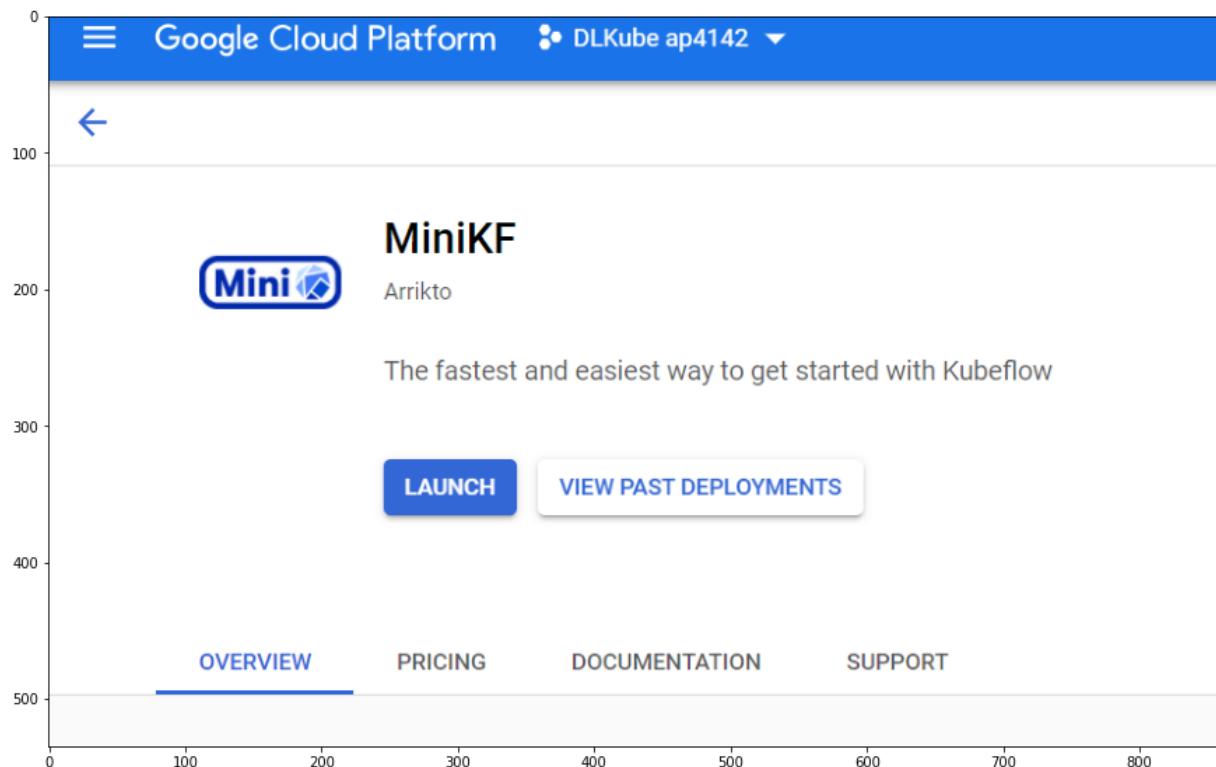
```
In [ ]: 1 from PIL import Image
2 import glob
3 import matplotlib.pyplot as plt
4
5 print('(a)')
6 files = glob.glob('3.1.1*')
7 for f in files:
8     fig = plt.gcf()
9     fig.set_size_inches(10, 7.5)
10    img = Image.open(f)
11    plt.imshow(img)
12    plt.show()
13
```

(a)



```
In [ ]:  
1 print('(b)')  
2 files = glob.glob('3.1.2*')  
3 for f in files:  
4     fig = plt.gcf()  
5     fig.set_size_inches(15,10)  
6     img = Image.open(f)  
7     plt.imshow(img)  
8     plt.show()
```

(b)



The screenshot shows a file browser interface with two main sections. The left section displays a tree view of files under 'minikf-ap4142': Overview - minikf-ap4142, minikf (selected), minikf-ap4142-firewall (firewall), minikf-ap4142-firewall-htp01 (firewall), minikf (vm\_instance.py), minikf-ap4142 vm instance, minikf-ap4142-data-disk (disk), and minikf-ap4142-password (password.py). A message at the top says 'minikf-ap4142 has been deployed'. The right section shows deployment details for 'minikf': MiniKF, Solution provided by Arrikto Inc., MiniKF dashboard (<https://minikf-ap4142.endpoints.dlhw5-334022.cloud.goog/>), MiniKF username (user), MiniKF password (k9Hy1kRxmq), Instance (minikf-ap4142), Instance zone (northamerica-northeast1-a), and Instance machine type (n1-standard-8). A 'MORE ABOUT THE SOFTWARE' link is also present.

The screenshot shows the 'New MiniKF deployment' configuration dialog. It includes sections for Data Disk (Data disk type: Standard Persistent Disk, size: 500 GB), GPUs (GPU settings), Networking (Network interfaces: default default (10.162.0.0/20), ADD NETWORK INTERFACE), and a 'MORE' section. At the bottom, a checkbox is checked for accepting terms of service, and a large blue 'DEPLOY' button is visible.

New MiniKF deployment

Data Disk

Data disk type: Standard Persistent Disk

Data disk size in GB: 500

GPUs

Networking

Network interfaces: default default (10.162.0.0/20)

ADD NETWORK INTERFACE

MORE

I accept the [GCP Marketplace Terms of Service](#) and [Arrikto Inc. Terms of Service](#).

DEPLOY

```
In [ ]:  
1 print('(c)')  
2 files = glob.glob('3.1.3*')  
3 for f in files:  
4     fig = plt.gcf()  
5     fig.set_size_inches(15,10)  
6     img = Image.open(f)  
7     plt.imshow(img)  
8     plt.show()
```

(c)

The screenshot shows the Kubeflow UI interface. On the left is a sidebar with a dark blue background containing various navigation options: Home, Notebooks, Tensorboards, Models, Snapshots, Volumes, Experiments (AutoML), Experiments (KFP), Pipelines, Runs, Recurring Runs, Artifacts, and Executions. The main area is titled "kubeflow-user (Owner)" and contains several sections:

- Quick shortcuts:** Includes links to "Upload a pipeline", "View all pipeline runs", "Create a new Notebook server", and "View Katib Experiments".
- Recent Notebooks:** A table showing "No Notebooks in namespace kubeflow-user".
- Recent Pipelines:** A table listing three pipelines:
  - [Tutorial] DSL - Control structures (Created 12/4/2021, 12:47:59 AM)
  - [Tutorial] Data passing in python components (Created 12/4/2021, 12:47:58 AM)
  - [Demo] TFX - Taxi tip prediction model trainer (Created 12/4/2021, 12:47:57 AM)
- Recent Pipeline Runs:** A table showing "None Found".
- Documentation:** A sidebar with links to "Getting Started with Kubeflow", "Minikube", "Micralk8s for Kubeflow", "Minikube for Kubeflow", "Kubeflow on GCP", "Kubeflow on AWS", and "Requirements for Kubeflow".

```
0 minikf@minikf-ap4142: ~ - Google Chrome
1 ssh.cloud.google.com/projects/dlhw5-334022/zones/northamerica-northeast1-a/instances/minikf-ap4142?authuser=2&hl=en_US&proj...
2 Connected, host fingerprint: ssh-rsa 0 BC:08:01:1F:B0:74:2E:8C:33:72:2C:68:86:7B
3 ECD52:E4:11:C0:56:05:A7:54:A4:75:D0:3B:4B:4E:97:FA:14
4 Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 5.4.104-0504104-generic x86_64)
5
6 * Documentation: https://help.ubuntu.com
7 * Management: https://landscape.canonical.com
8 * Support: https://ubuntu.com/advantage
9
10 126 packages can be updated.
11 74 updates are security updates.
12
13
14 The programs included with the Ubuntu system are free software;
15 the exact distribution terms for each program are described in the
16 individual files in /usr/share/doc/*copyright.
17
18 Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
19 applicable law.
20
21 Switching you to user "minikf"...
22
23 Welcome to MiniKF!
24 Type "minikf" to ensure everything is up and running.
25 minikf@minikf-ap4142:~$ minikf
```

The screenshot shows the Kubeflow UI interface. On the left, there is a sidebar with the following menu items:

- Home
- Notebooks
- Tensorboards
- Models
- Snapshots
- Volumes
- Experiments (AutoML)
- Experiments (KFP)

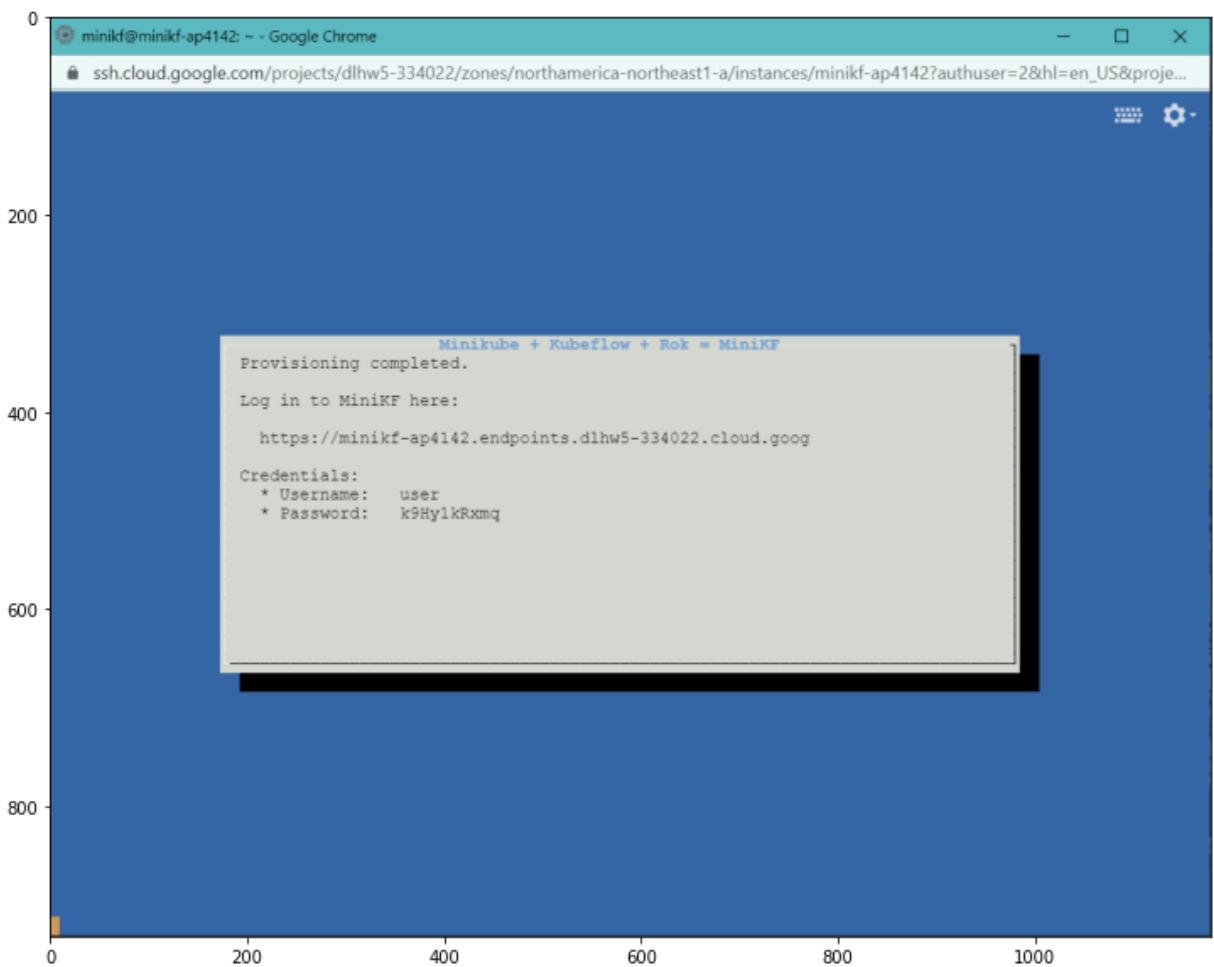
On the right, the main content area displays the following message:

kubeflow-user (Owner) ▾

Rok

+ New bucket

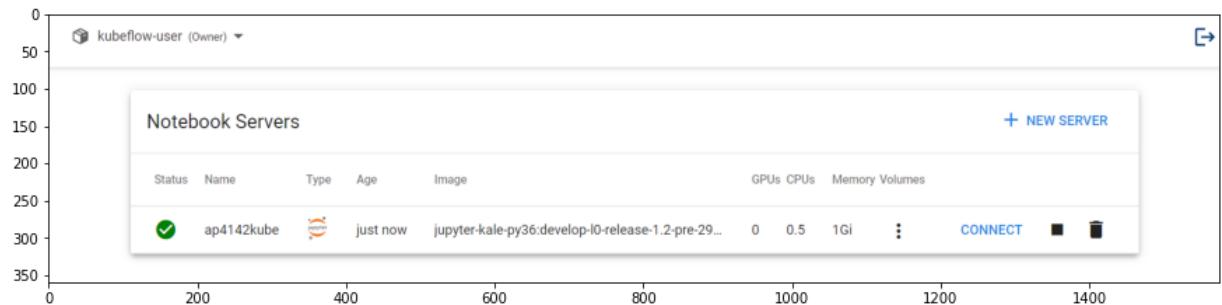
No buckets yet  
Add one!



In [ ]:

```
1 print('(a)')
2 files = glob.glob('3.2.1*')
3 for f in files:
4     fig = plt.gcf()
5     fig.set_size_inches(15,10)
6     img = Image.open(f)
7     plt.imshow(img)
8     plt.show()
```

(a)

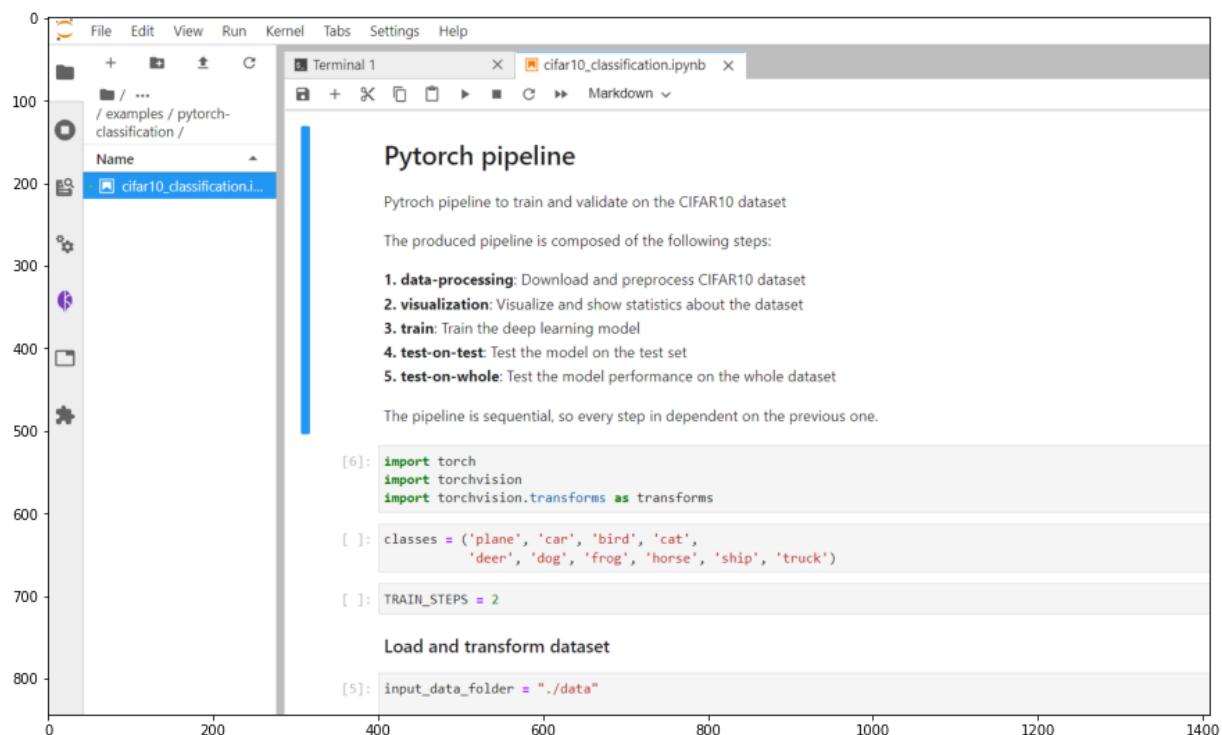
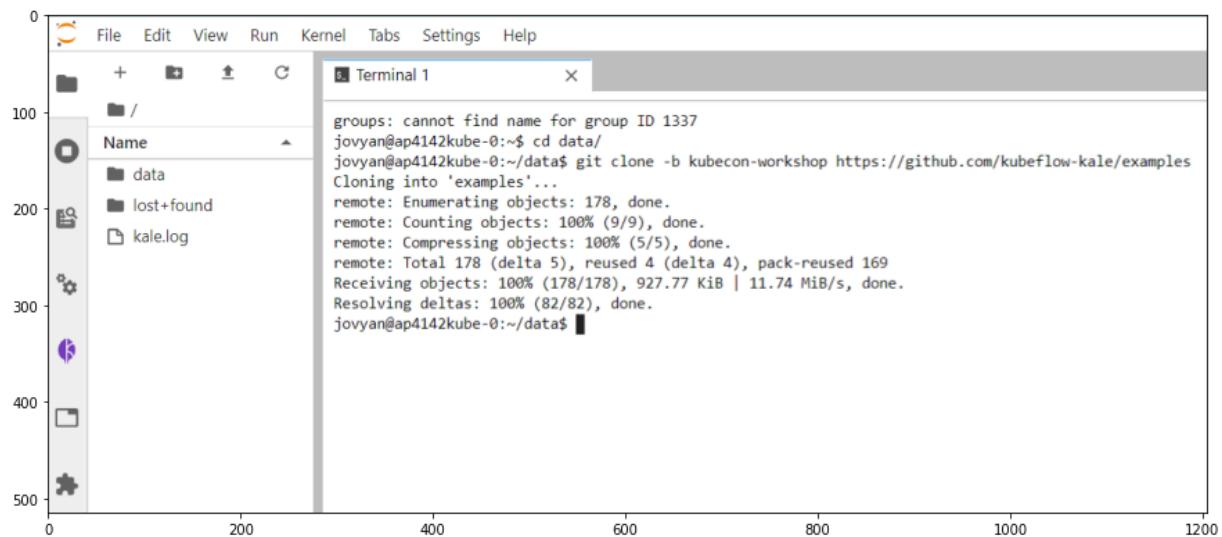


The screenshot shows the Kubeflow Notebook configuration dialog. On the left is a sidebar with links: Home, Notebooks (selected), Tensorboards, Models, Snapshots, Volumes, Experiments (AutoML), Experiments (KFP), Pipelines, Runs, Recurring Runs, Artifacts, and Executions. The main area has tabs for "Variables.", "Configurations", and "Environment". Under "Configurations", there is a note: "Allow access to Kubeflow Pipelines, Allow access to Rok". Under "Environment", there is a note: "None". Below these are sections for "Affinity / Tolerations" and "Miscellaneous Settings". In "Affinity / Tolerations", there is a note: "Configure the Notebook's Affinity and Tolerations." and dropdowns for "Affinity Config" (set to "None") and "Tolerations Group" (set to "None"). In "Miscellaneous Settings", there is a note: "Other possible settings to be applied to the Notebook Server." and a checked checkbox for "Enable Shared Memory". At the bottom are "LAUNCH" and "CANCEL" buttons.

The screenshot shows the Kubeflow UI interface. On the left, there is a sidebar with a dark blue background containing various navigation options: Home, Notebooks (selected), Tensorboards, Models, Snapshots, Volumes, Experiments (AutoML), Experiments (KFP), Pipelines, Runs, Recurring Runs, Artifacts, and Executions. The main area has a light gray background. At the top right, it says "kubeflow-user (Owner)". Below that, there is a section titled "Rok JupyterLab URL" with a placeholder "Load an existing Jupyter Lab by providing a valid Rok URL." A search icon is located at the top right of this section. Below this, there is a "Name" input field with "ap4142kube" typed in, and a "Namespace" input field with "kubeflow-user". Underneath these fields, there is a "Image" section with a sub-section "A started Jupyter Docker Image with a baseline deployment and typical ML packages". It includes a checkbox for "Custom Image" and three pre-defined image options: "jupyterlab" (selected), "Visual Studio Code", and "R Studio". The "jupyterlab" option has the URL "gcr.io/arrikto/jupyter-kale-py36:develop-l0-release-1.2-pre-295-g622fe91aca" listed below it.

```
In [ ]:
1 print('(b)')
2 files = glob.glob('3.2.2*')
3 for f in files:
4     fig = plt.gcf()
5     fig.set_size_inches(15,10)
6     img = Image.open(f)
7     plt.imshow(img)
8     plt.show()
```

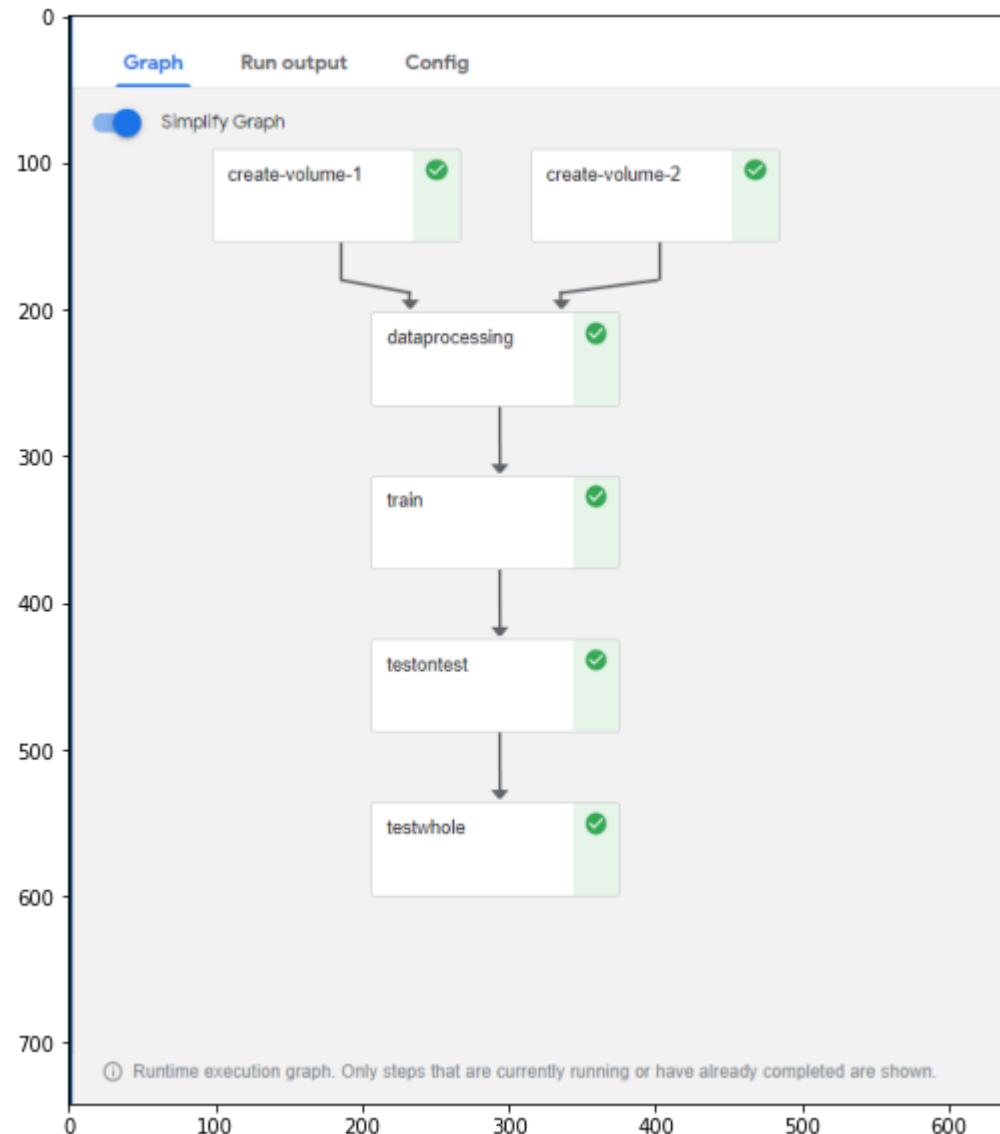
(b)





```
In [ ]:  
1 print('(c)')  
2 files = glob.glob('3.2.3*')  
3 c=0  
4 for f in files:  
5     if(c<4):  
6         fig = plt.gcf()  
7         fig.set_size_inches(15,10)  
8         img = Image.open(f)  
9         plt.imshow(img)  
10        plt.show()  
11    c+=1
```

(c)





The screenshot shows the Kale Deployment Panel interface. On the left, there's a sidebar with sections for Pipeline Metadata (Experiment Name: cifar10, Pipeline Name: cifar10-classification, Pipeline Description: Sequential PyTorch pipeline to train a), Run (HP Tuning with Katib), and Advanced Settings. A large blue button at the bottom says "COMPILE AND RUN". On the right, a terminal window titled "Terminal 1" shows code for a PyTorch pipeline to train and validate on the CIFAR10 dataset. The code includes imports for torch, torchvision, and transforms, a list of classes, and a parameter TRAIN\_STEPS set to 2. The terminal also shows a progress bar for saving completed.

```
File Edit View Run Kernel Tabs Settings Help
Kale Deployment Panel
Enable
Pipeline Metadata
Select experiment + New Experiment
Experiment Name: cifar10
Pipeline Name: cifar10-classification
Pipeline Description: Sequential PyTorch pipeline to train a
Run
HP Tuning with Katib
SET UP KATIB JOB
Advanced Settings
COMPILE AND RUN
```

```
Pytorch pipeline 1
Pytorch pipeline to train and validate on the CIFAR10 dataset
The produced pipeline is composed of the following steps:
1. data-processing: Download and preprocess CIFAR10 dataset
2. visualization: Visualize and show statistics about the dataset
3. train: Train the deep learning model
4. test-on-test: Test the model on the test set
5. test-on-whole: Test the model performance on the whole dataset
The pipeline is sequential, so every step is dependent on the previous one.

imports
[1]: import torch
import torchvision
import torchvision.transforms as transforms

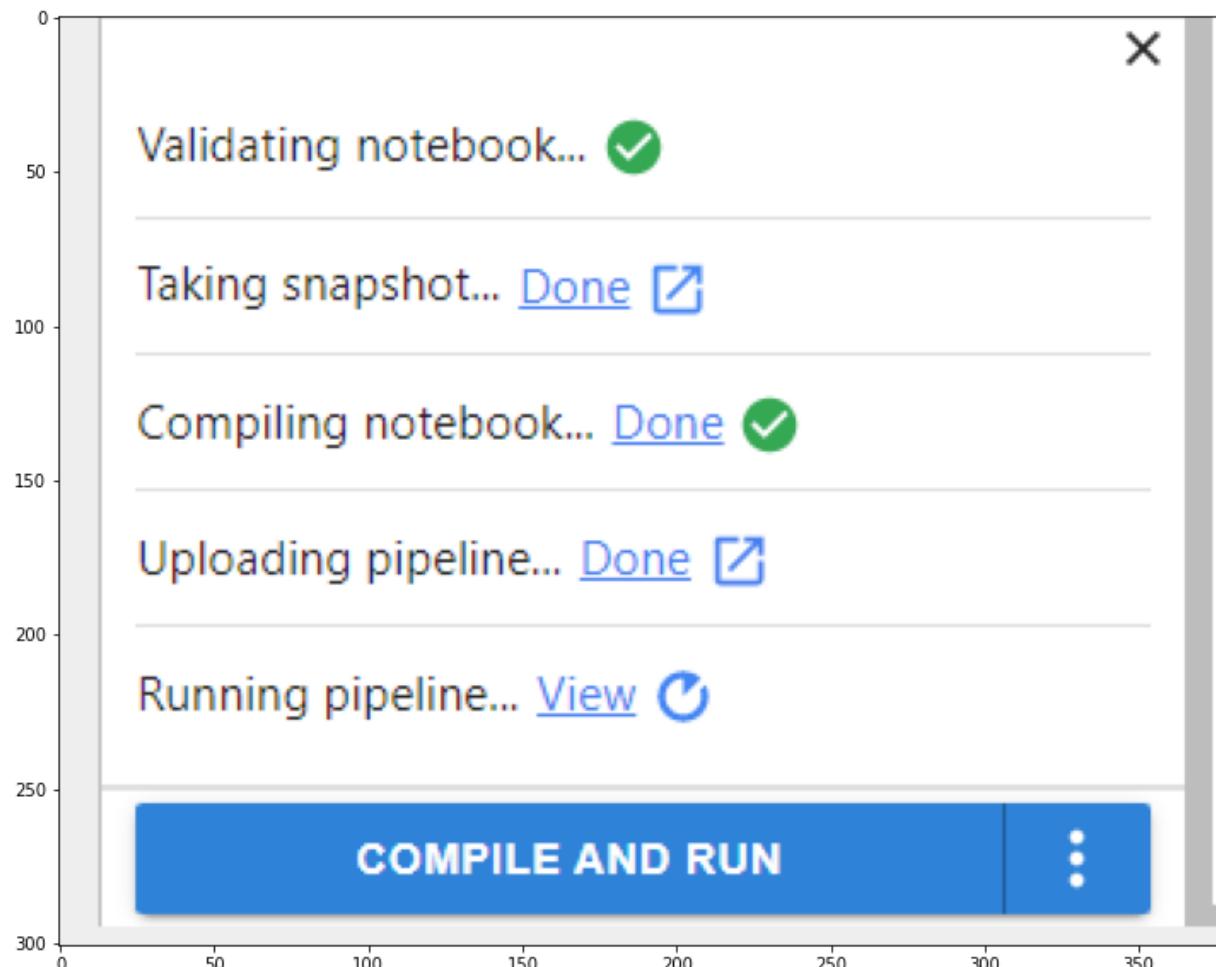
functions
[2]: classes = ('plane', 'car', 'bird', 'cat',
              'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

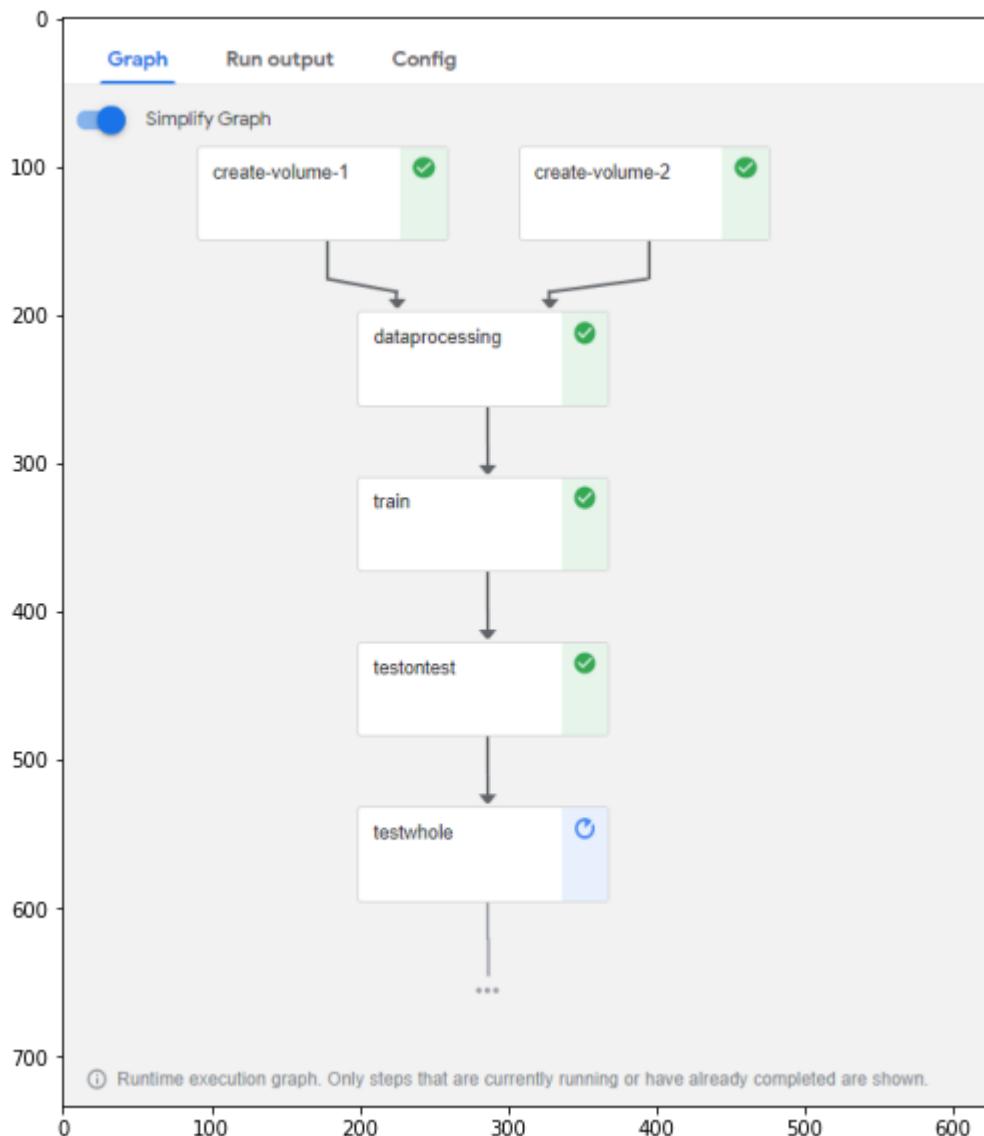
pipeline-parameters
[3]: TRAIN_STEPS = 2

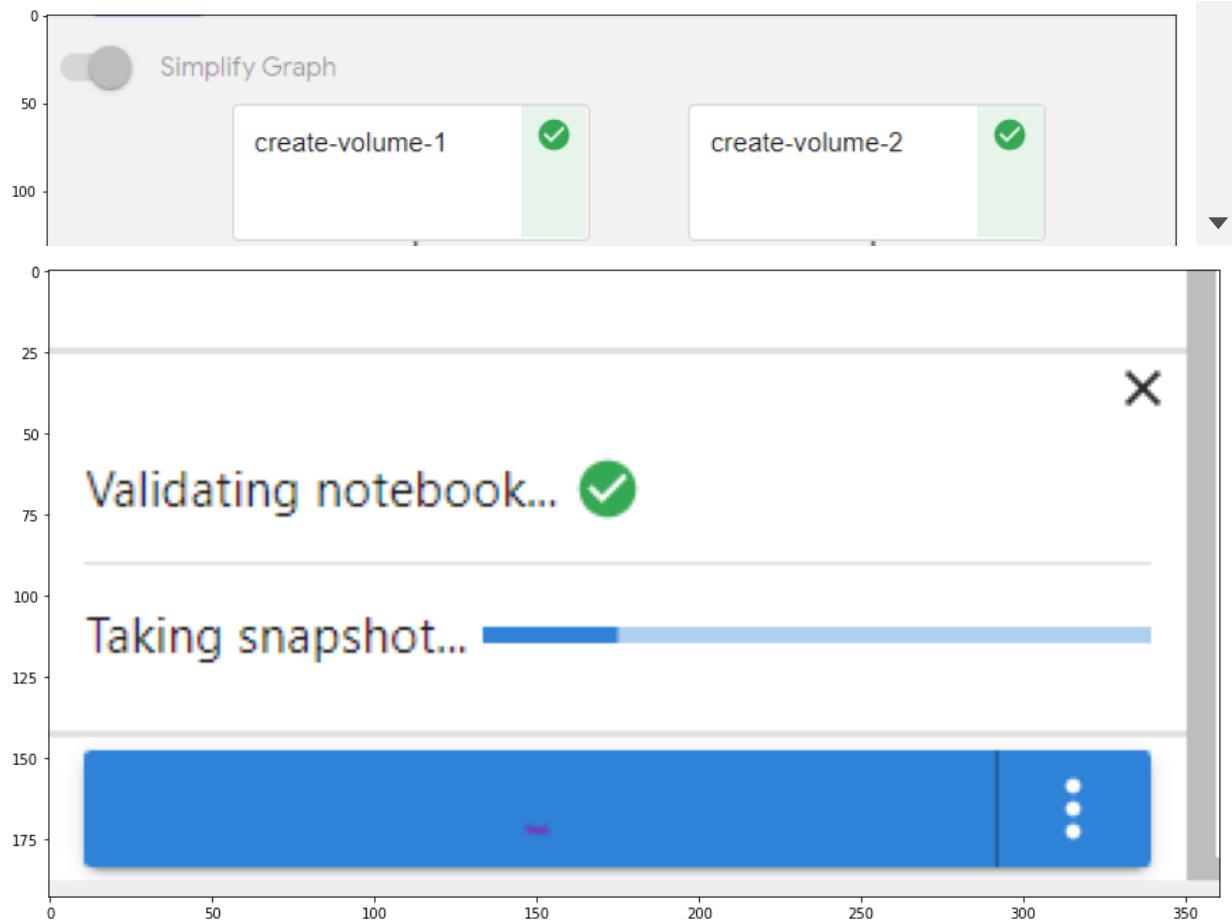
Load and transform dataset
step: dataprocessing
```

```
In [ ]:  
1 print('(c)')  
2 files = glob.glob('3.2.3*')  
3 c=0  
4 for f in files:  
5     if(c>=4):  
6         fig = plt.gcf()  
7         fig.set_size_inches(15,10)  
8         img = Image.open(f)  
9         plt.imshow(img)  
10        plt.show()  
11    c+=1
```

(c)

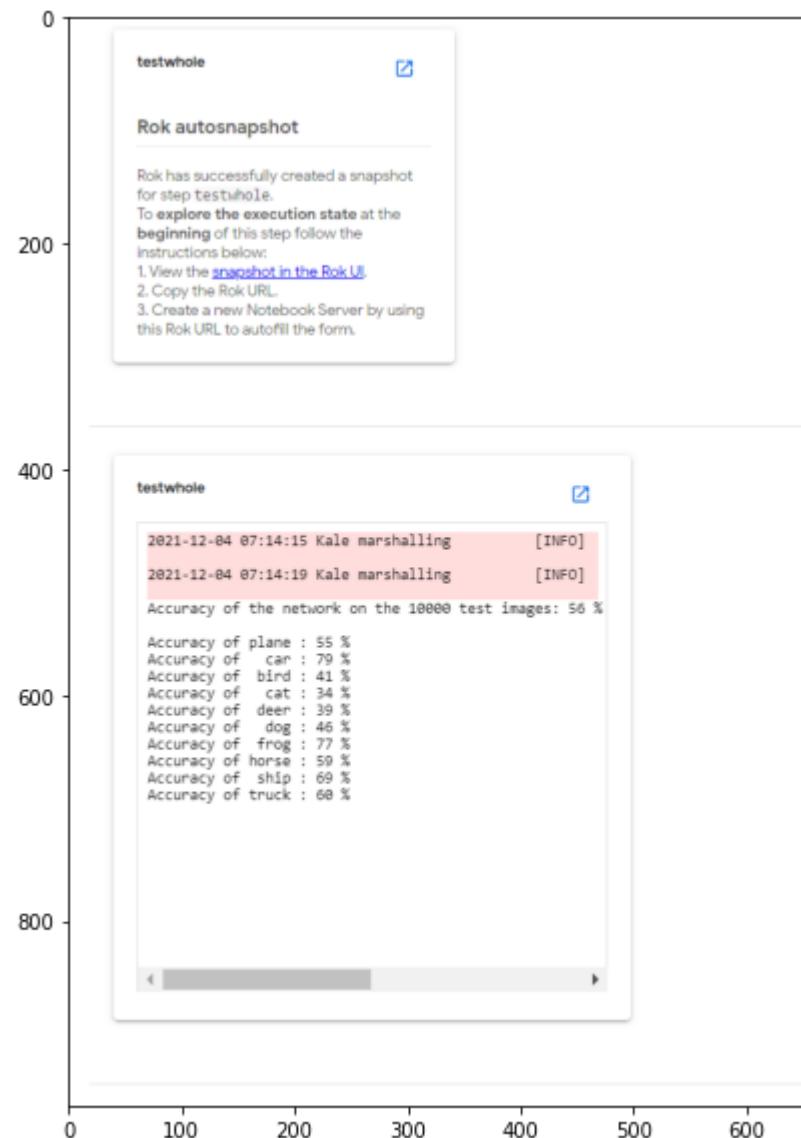


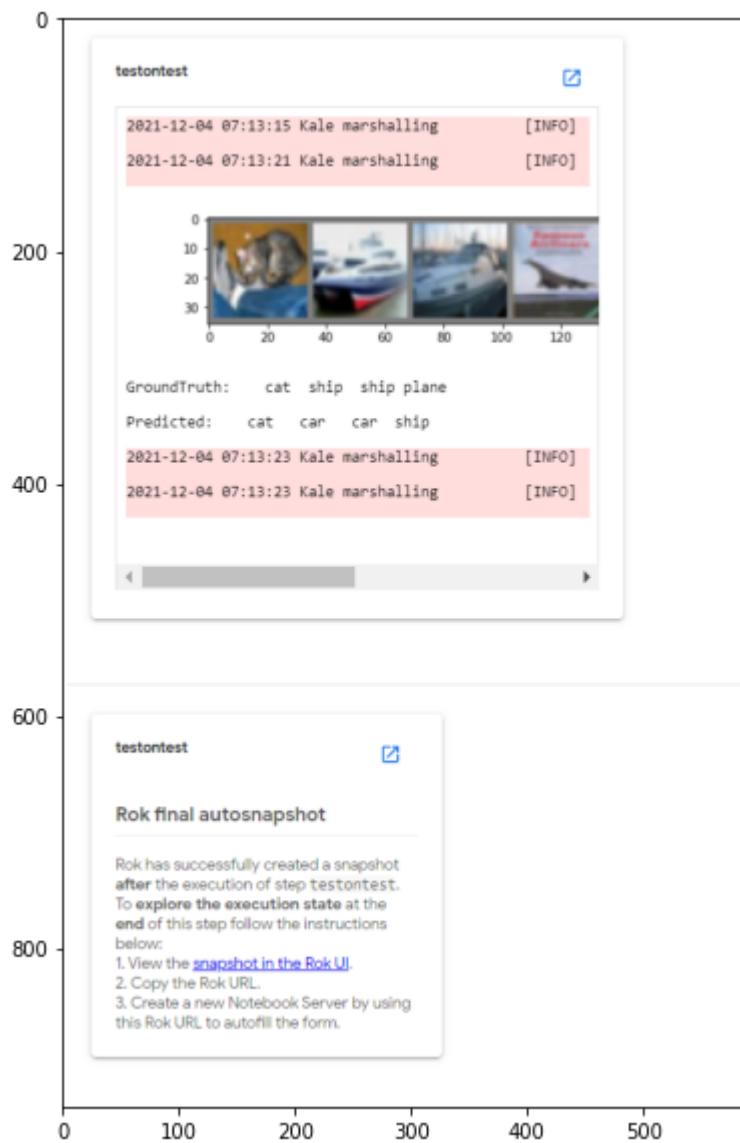


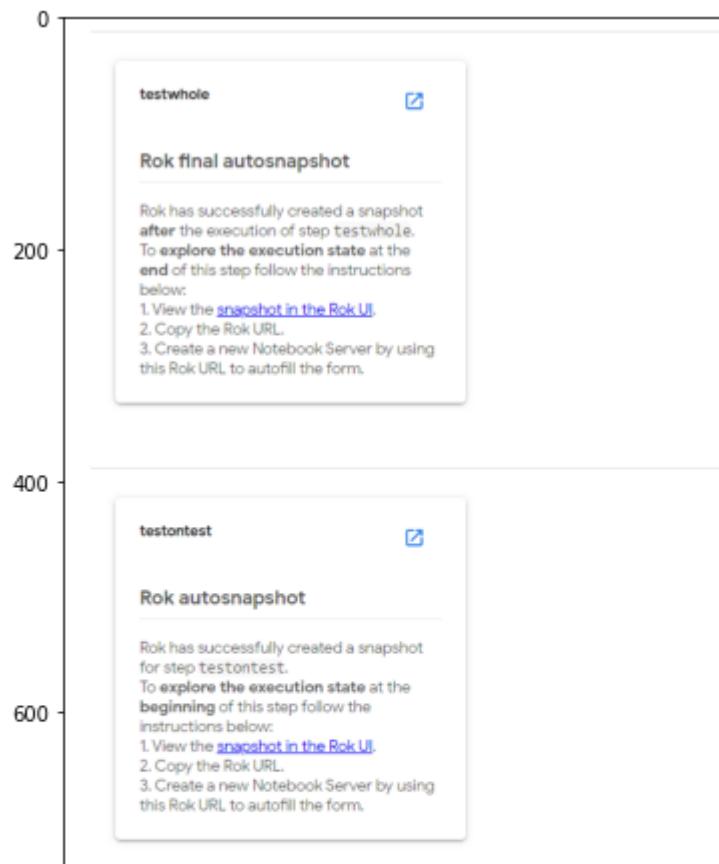


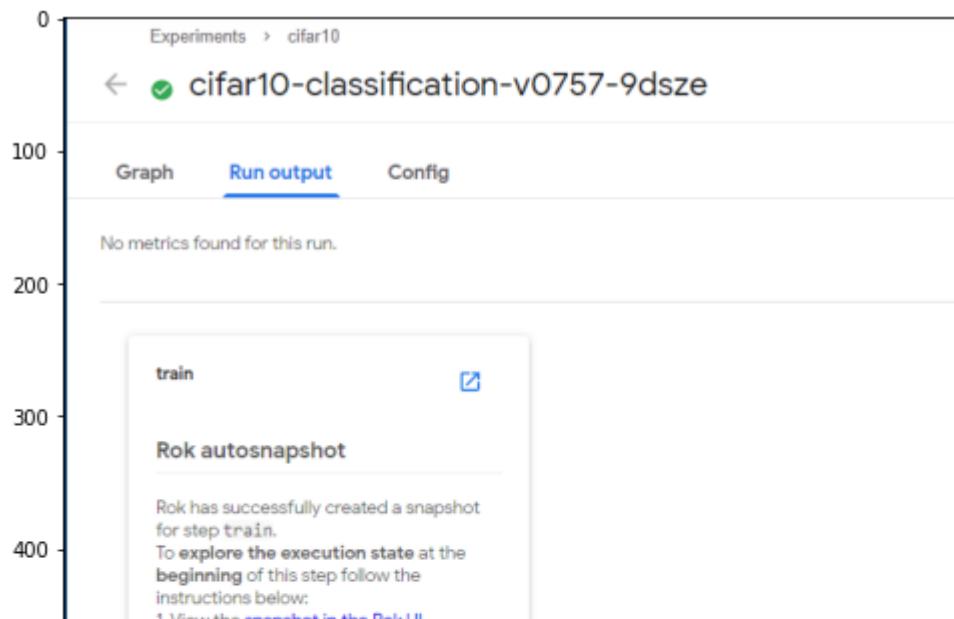
```
In [ ]: 1 print('(d)')
2 files = glob.glob('3.2.4*')
3 c=0
4 for f in files:
5     if(c<4):
6         fig = plt.gcf()
7         fig.set_size_inches(15, 10)
8         img = Image.open(f)
9         plt.imshow(img)
10        plt.show()
11    c+=1
```

(d)



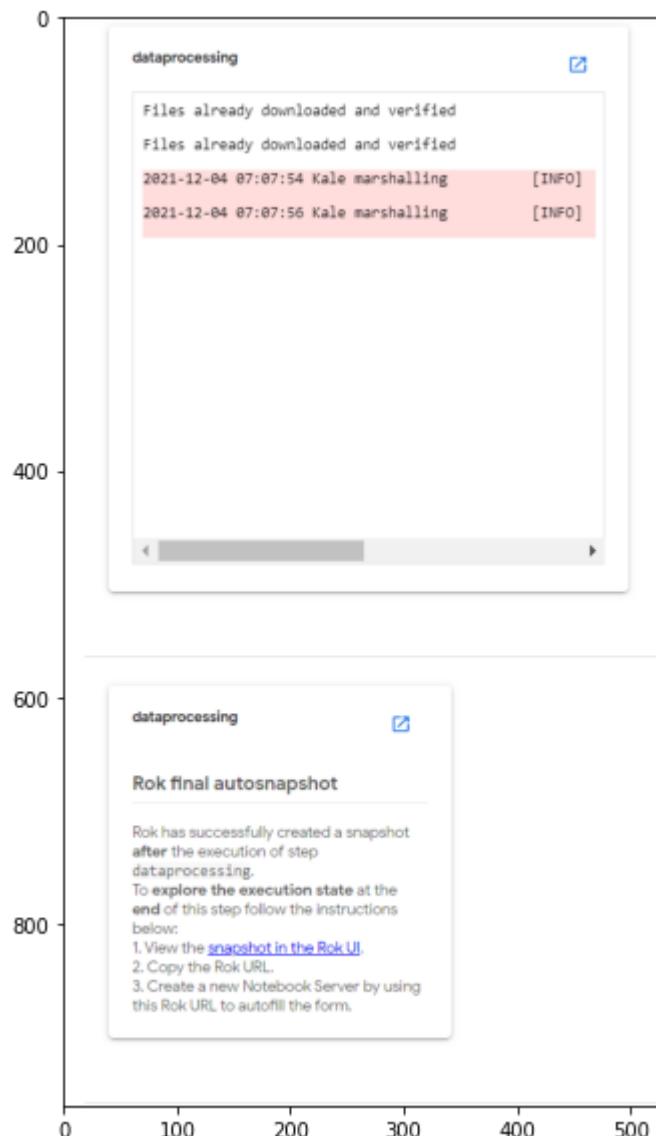


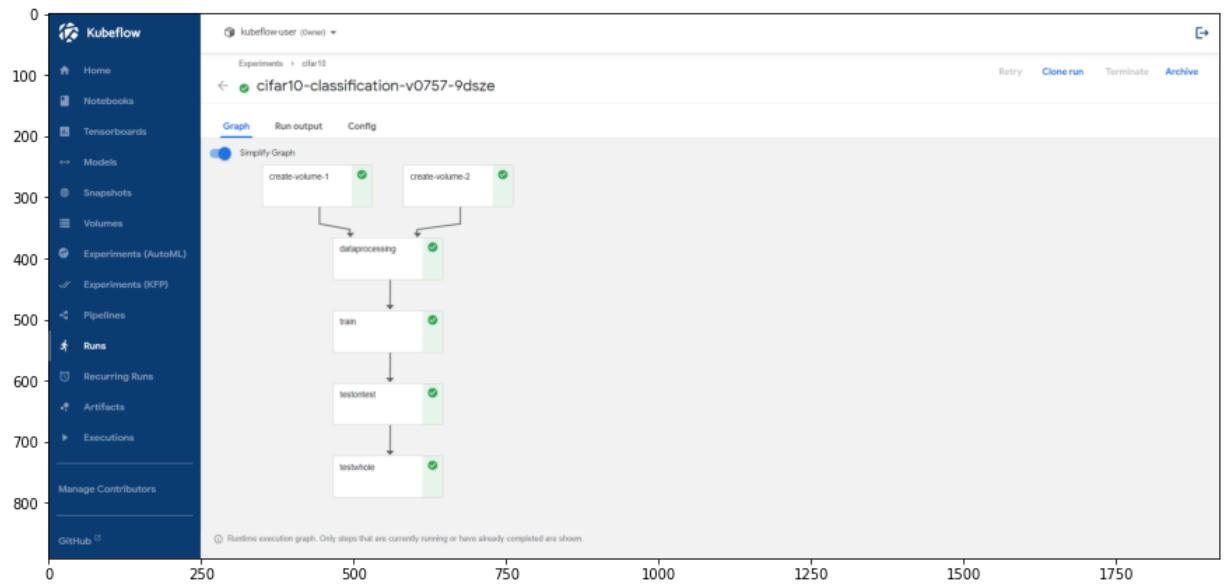




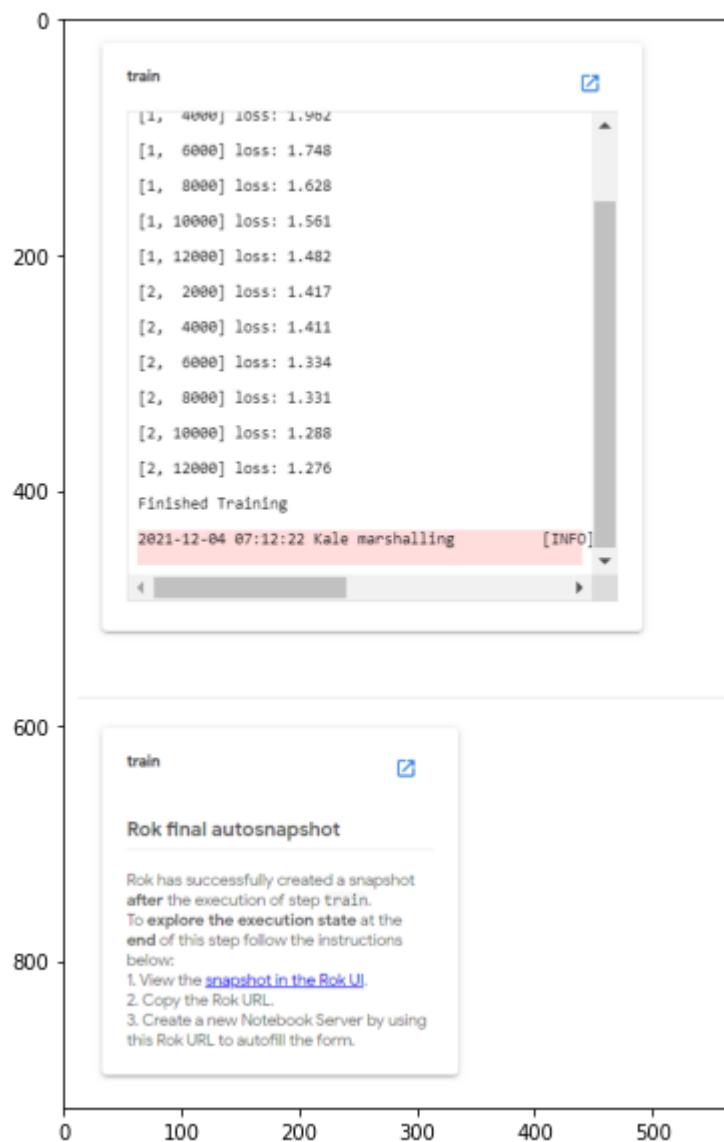
```
In [ ]: 1 print('(d)')
2 files = glob.glob('3.2.4*')
3 c=0
4 for f in files:
5     if(c>=4):
6         fig = plt.gcf()
7         fig.set_size_inches(15,10)
8         img = Image.open(f)
9         plt.imshow(img)
10        plt.show()
11        c+=1
```

(d)



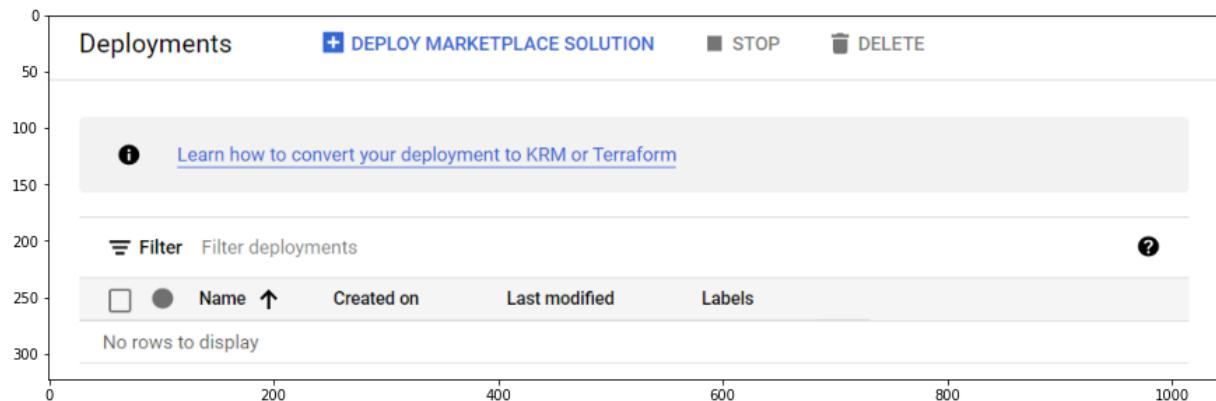
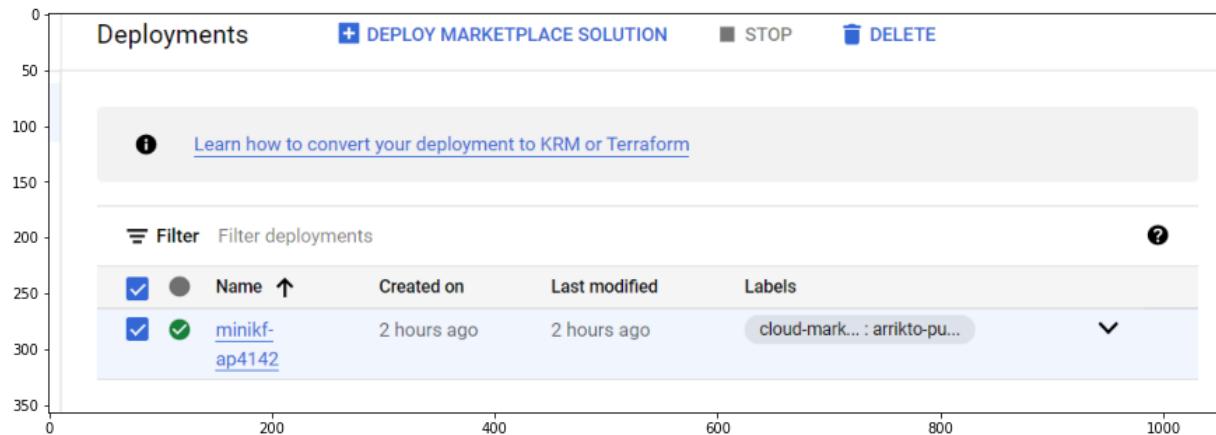






```
In [ ]: 1 print('(e)')
2 files = glob.glob('3.2.5*')
3 c=0
4 for f in files[:-1]:
5     fig = plt.gcf()
6     fig.set_size_inches(15,10)
7     img = Image.open(f)
8     plt.imshow(img)
9     plt.show()
```

(e)



```
In [ ]: 1
```

## Problem 4

1. An episodic task is a task which has an end state or a goal after which the episode terminates that is it has at least one finite state. An example of episodic task is a game like tic-tac-toe in which an episode is finished after the game terminates and a new one starts.  
A continuous task is a task which does not terminate. For example prediction of stock market for trading. This task has no end to it therefore it is a continuous task.
2. In reinforcement learning Exploration corresponds to the case of choosing an action randomly(explored or unexplored) which may not maximize the reward from the information we have at that point. Whereas, exploitation corresponds to the case of greedily choosing the action which will maximize the reward even if there are unexplored actions.

$\epsilon$ -greed method is used to make sure that the model does not always make the greed choice and explores all the possible actions in the action space instead of always choosing only one path.

A decaying  $\epsilon$  will perform better than a fixed  $\epsilon$ . So following a schedule is more preferable than having a fixed  $\epsilon$ . This is because in the first few iterations, we will be able to cover the action space after which choosing actions more greedily is preferred.

Having  $\epsilon = 0$  will ensure that actions are picked greedily so certain explored actions will be chosen over unexplored actions which could give better results. So choosing a random action with a non-zero probability of  $\epsilon$  will make sure that the unexplored actions are also selected and are explored. The higher the  $\epsilon$  more the exploration. So a decaying  $\epsilon$  will make sure we do enough exploration in the beginning and follow exploitation towards the end.

3. In Q-learning the values of each state-action pairs is stored in a table format. But in deep Q-learning we use a neural network to learn the Q function.

Deep Q-learning uses experience relay to store the experience tuple after each time step.

Deep Q-learning has 2 separate networks, a Q-network and a target Q-network.

After initializing the network we first choose the action using an  $\epsilon$ -greed policy. We then apply the action chosen to observe the reward from the environment and the new state, and generate an experience tuple which is stored in the experience buffer D. Then we select a random minibatch of the tuples stored in D to train the Q-network. The loss function is given as  $(y_j - Q(\phi_j, a_j; \theta))^2$  where  $y_j = r_j$  if the new state is terminal and

$$y_j = r_j + \max_{a'} Q(\phi_{j+1}, a'; \theta) \text{ otherwise.}$$

4. In the case of Deep Q-learning, Q values are generated using a neural network. So a change in the network changes the Q value for all the actions. A target network is required to handle the moving target problem as the update equation for the Q value depends on the Q value of another action as well. It helps the network achieve more stability as updating the network at each time step may affect the stability and convergence of the model. It makes training more stable because we use a minibatch of data to update the network instead of doing it at every time step.
5. Updating the network after each iteration can cause a lot of instability in the model. Using a minibatch of experience tuples from the experience buffer D helps training to be more stable. Another advantage is that as we randomly choose a minibatch from D the tuples will be trained multiple times resulting in more efficient learning. As some of the experience tuples come from previous versions of the Q network, there is also less overfitting.
6. Prioritized experience replay is used to improve data efficiency. In deep RL it is used in replay buffers which allows for more efficient learning by replaying important transitions more frequently. This is because not all experiences are equal and some experiences contribute more to the agent's learning than others. Generally more importance is given to the most 'surprising' experiences.

## 7. Similarities:

- 1) Both GORILA and Ape-X architectures follow 2 parts - acting & learning.
- 2) Both GORILA and Ape-X run with no high level synchronization.
- 3) In both GORILA and Ape-X acting and learning run concurrently.
- 4) For both, acting & learning could be distributed across multiple workers.

## Differences:

- 1) Ape-X uses a shared, centralized replay memory unlike Gorila
- 2) Gorila uses uniform sampling, whereas Ape-X prioritizes the most useful data.
- 3) Since priorities are shared, high priority data discovered by any actor can benefit the whole system which does not happen in Gorila.

In [ ]:

1